



User Guide

# Microsoft SQL Server on Amazon EC2



# Microsoft SQL Server on Amazon EC2: User Guide

Copyright © 2026 Amazon Web Services, Inc. and/or its affiliates. All rights reserved.

Amazon's trademarks and trade dress may not be used in connection with any product or service that is not Amazon's, in any manner that is likely to cause confusion among customers, or in any manner that disparages or discredits Amazon. All other trademarks not owned by Amazon are the property of their respective owners, who may or may not be affiliated with, connected to, or sponsored by Amazon.

# Table of Contents

|                                                               |           |
|---------------------------------------------------------------|-----------|
| <b>What is Microsoft SQL Server on Amazon EC2?</b>            | <b>1</b>  |
| Features                                                      | 1         |
| Pricing                                                       | 3         |
| Options to run SQL Server                                     | 3         |
| SQL Server on Amazon EC2                                      | 4         |
| RDS for SQL Server                                            | 4         |
| Amazon RDS Custom                                             | 5         |
| Concepts and terminology                                      | 6         |
| SQL Clustering best practices                                 | 9         |
| Assign IP addresses                                           | 10        |
| Cluster properties                                            | 11        |
| Cluster quorum votes and 50/50 splits in a multi-site cluster | 11        |
| DNS registration                                              | 12        |
| Elastic Network Adapters (ENAs)                               | 13        |
| Multi-site clusters and EC2 instance placement                | 13        |
| Instance type selection                                       | 13        |
| Assign elastic network interfaces and IPs to the instance     | 13        |
| Heartbeat network                                             | 14        |
| Configure the network adapter in the OS                       | 14        |
| IPv6                                                          | 14        |
| Host record TTL for SQL Availability Group Listeners          | 14        |
| Logging                                                       | 15        |
| NetBIOS over TCP                                              | 15        |
| NetFT Virtual Adapter                                         | 15        |
| Set possible owners                                           | 16        |
| Tune the failover thresholds                                  | 17        |
| Witness importance and Dynamic Quorum Architecture            | 18        |
| Troubleshoot                                                  | 18        |
| <b>Set up SQL Server on Amazon EC2</b>                        | <b>20</b> |
| Prerequisites                                                 | 20        |
| Sign up for an AWS account                                    | 20        |
| Create a key pair                                             | 20        |
| Create a security group                                       | 21        |
| Permissions                                                   | 24        |

|                                                                 |           |
|-----------------------------------------------------------------|-----------|
| <b>Licensing options and considerations .....</b>               | <b>25</b> |
| Licensing options .....                                         | 25        |
| License-included .....                                          | 25        |
| BYOL .....                                                      | 26        |
| Licensing considerations .....                                  | 26        |
| Choose a SQL Server edition .....                               | 27        |
| Purchase SQL Server from AWS .....                              | 27        |
| Use BYOL for SQL Server on AWS .....                            | 28        |
| Quantify license requirements .....                             | 28        |
| License Mobility with SQL Server .....                          | 28        |
| Track BYOL license consumption .....                            | 29        |
| SQL Server CALs .....                                           | 29        |
| Licensing for passive failover .....                            | 29        |
| Amazon EC2 High Availability for SQL Server on Amazon EC2 ..... | 30        |
| Supported deployments .....                                     | 30        |
| Considerations and requirements .....                           | 31        |
| Prerequisites .....                                             | 31        |
| How it works .....                                              | 32        |
| Getting started .....                                           | 33        |
| Disable Amazon EC2 High Availability for SQL Server .....       | 43        |
| View states .....                                               | 44        |
| <b>Find a SQL Server license-included AMI .....</b>             | <b>45</b> |
| Methods to find a SQL Server license-included AMI .....         | 45        |
| <b>Deploy SQL Server on Amazon EC2 .....</b>                    | <b>49</b> |
| Considerations .....                                            | 49        |
| Deployment options .....                                        | 50        |
| <b>Connect to SQL Server on Amazon EC2 .....</b>                | <b>58</b> |
| SSMS .....                                                      | 58        |
| Configuration Manager .....                                     | 59        |
| <b>VSS based database backup .....</b>                          | <b>60</b> |
| Snapshot prerequisites .....                                    | 60        |
| Create AWS VSS solution based EBS snapshots .....               | 61        |
| <b>VSS based database restore with automation runbook .....</b> | <b>62</b> |
| Pricing .....                                                   | 62        |
| VSS based database restore solution change history .....        | 63        |
| Prerequisites .....                                             | 63        |

|                                                                                   |            |
|-----------------------------------------------------------------------------------|------------|
| Grant IAM permissions for the restore process .....                               | 65         |
| Restore from VSS snapshots .....                                                  | 69         |
| Runbook parameters .....                                                          | 69         |
| Run the restore process .....                                                     | 71         |
| Troubleshoot restoring your SQL Server database using AWS VSS solution .....      | 74         |
| <b>Evaluate downgrading your SQL Server edition .....</b>                         | <b>76</b>  |
| Downgrade requirements .....                                                      | 76         |
| Downgrade your SQL Server Enterprise edition .....                                | 78         |
| <b>Migrating an on-premises database to Amazon EC2 .....</b>                      | <b>80</b>  |
| Automated SQL Server backup and restore .....                                     | 80         |
| Manual SQL Server backup and restore .....                                        | 80         |
| Prerequisites .....                                                               | 81         |
| Step 1: Backing up your database .....                                            | 81         |
| Step 2: Uploading your database backup files .....                                | 81         |
| Step 3: Downloading your database backup files .....                              | 82         |
| Step 4: Restoring your database backup files .....                                | 82         |
| Server rehost .....                                                               | 82         |
| <b>Migrate Microsoft SQL Server from Windows to Linux .....</b>                   | <b>84</b>  |
| Concepts .....                                                                    | 84         |
| Related services .....                                                            | 85         |
| How Windows to Linux replatforming assistant for Microsoft SQL Server works ..... | 85         |
| Components .....                                                                  | 86         |
| Replatforming script prerequisites .....                                          | 86         |
| Prerequisites to run the replatforming script .....                               | 86         |
| Prerequisites for replatforming to an existing EC2 instance .....                 | 88         |
| Run the replatforming script .....                                                | 88         |
| Replatforming script examples .....                                               | 89         |
| Replatforming script parameters .....                                             | 91         |
| <b>Security .....</b>                                                             | <b>95</b>  |
| Identity and access management .....                                              | 95         |
| AWS managed policies .....                                                        | 96         |
| Service-linked role .....                                                         | 99         |
| <b>Document history .....</b>                                                     | <b>101</b> |

# What is Microsoft SQL Server on Amazon EC2?

You can run Microsoft SQL Server on Amazon EC2. Microsoft SQL Server is a relational database management system (RDBMS) whose primary purpose is to store and retrieve data. SQL Server includes additional services, such as Analysis Services (SSAS), Reporting Services (SSRS), Integration Services (SSIS), and Machine Learning (ML). AWS provides a comprehensive set of services and tools to deploy Microsoft SQL Server on the reliable and secure AWS Cloud infrastructure. The benefits of running SQL Server on AWS include cost savings, scalability, high availability and disaster recovery, improved performance, and ease of management. For more information, see [Learn why AWS is the best cloud to run Microsoft Windows Server and SQL Server workloads](#) on the AWS Compute blog.

Amazon Elastic Compute Cloud (Amazon EC2) supports a self-managed SQL Server. That is, it gives you full control over the setup of the infrastructure and the database environment. Running SQL Server on Amazon EC2 is very similar to running SQL Server on your own server. You have full control of the database and operating system-level access, so you can use your choice of tools to manage the operating system, database software, patches, data replication, backup, and restoration. You are responsible for data replication and recovery across your instances in the same or different AWS Regions. For more information, refer to the [AWS Shared Responsibility Model](#).

## Overview topics

- [Microsoft SQL Server on Amazon EC2 features](#)
- [Microsoft SQL Server on Amazon EC2 pricing](#)
- [Options to run SQL Server on the AWS Cloud](#)
- [Microsoft SQL Server on Amazon EC2 concepts and terminology](#)
- [Best practices and recommendations for SQL Server clustering on Amazon EC2](#)

## Microsoft SQL Server on Amazon EC2 features

SQL Server on Amazon EC2 provides the following features:

- **Flexible licensing options** — When you use Amazon EC2 instances with the license included, you are using instances with fully-compliant Windows Server and SQL Server that are licensed through AWS. Flexible BYOL options include default tenant EC2 for products that are eligible for [Microsoft License Mobility through Software Assurance](#), as well as [Amazon EC2 Dedicated Hosts](#)

and [Amazon EC2 Dedicated Instances](#). You can use [AWS License Manager](#) to track the usage of software licenses and reduce the risk of non-compliance. For more information, see [Licensing](#) in the *Amazon Web Services and Microsoft Frequently Asked Questions*.

- **High performance block storage** — [Amazon Elastic Block Store](#) provides multiple options for high-performance block storage for Microsoft SQL Server. EC2 Instances using [io2 Block Express](#) give you the highest block storage performance with a single storage volume. Other SSD-backed Amazon EBS options include io2 volumes for business-critical applications and gp3 volumes for general purpose applications. Amazon EBS also offers crash-consistent snapshots, and enables application-consistent snapshots through Windows VSS (Volume Shadow Copy Services) to help protect your SQL Server deployments.
- **Fully-managed shared storage** — [Amazon FSx for Windows File Server](#) and Amazon FSx for NetApp ONTAP offer fully-managed shared storage for high-availability SQL Server failover cluster instances (FCI) workloads.
- **Windows-based services** — [Directory Service](#) offers managed Microsoft Active Directory with identity and access management.
- **Scalable processors** — [Intel Xeon Scalable Processors on AWS](#) provide you with better data protection, faster processing of more data volumes, and increased service flexibility for Amazon EC2.
- **Migration programs** — AWS offers programs for migration for customers looking to migrate SQL Server workloads to AWS. AWS [Migration Acceleration Program \(MAP\) for Windows](#) provides services, best practices, and tools to help you save costs and accelerate your migration on AWS.
- **Windows workload optimization** — After you move your SQL Server workloads to AWS, you can continue to optimize costs, usage, and licenses to suit your business requirements. With [Cost Explorer Service](#), you can visualize, understand, and manage your AWS costs and usage over time. [AWS Compute Optimizer](#) recommends optimal AWS compute resources for your workloads so that you can reduce costs up to 25% by analyzing historical utilization data. [AWS Trusted Advisor](#) can check that your EC2 instances have the required amount of SQL Server licenses and that the EC2 instance vCPU count doesn't exceed what is permitted for the SQL Server edition. [AWS Managed Services](#) can help operate your cloud environment post-migration by analyzing alerts and responding to incidents, reducing operational overhead and risk. You can use [AWS Systems Manager](#) to automate operational tasks across your AWS resources and better manage your infrastructure at scale.

AWS can help you to modernize you Windows-based applications with AWS open source services if you want to reduce the high cost of commercial licensing. Options include running SQL Server database applications on Linux, moving workloads to [Amazon Aurora](#), containerizing

your Windows applications with [Amazon EKS](#), going serverless with [AWS Lambda](#), or taking advantage of micro-services based architecture.

For more features specific to Amazon EC2, see [Features of Amazon EC2](#).

## Microsoft SQL Server on Amazon EC2 pricing

For information about pricing for Amazon EC2, see the [Amazon EC2 pricing](#) page.

For information about creating a price estimate for Microsoft Windows Server and Microsoft SQL Server, see [Tutorial: Using Windows Server and SQL Server on Amazon EC2 calculator](#) in the *AWS Pricing Calculator User Guide*.

## Options to run SQL Server on the AWS Cloud

AWS provides the option to run Microsoft SQL Server in a cloud environment. For developers and database administrators, running SQL Server in the AWS Cloud is similar to running SQL Server databases in a data center. There are three primary options to run SQL Server on AWS:

- Microsoft SQL Server on Amazon EC2
- Amazon RDS for Microsoft SQL Server
- Amazon RDS Custom for Microsoft SQL Server

Your application requirements, database features, functionality, growth capacity, and overall architecture complexity will determine which option to choose. Many AWS customers run multiple SQL Server database workloads across Amazon RDS and Amazon EC2. For more information on how to choose how to run SQL Server on the AWS Cloud, see [Decision matrix](#) on the AWS Prescriptive Guidance website.

If you are migrating multiple SQL Server databases to AWS, some of them might be a great fit for Amazon RDS, whereas others might be better suited to run directly on Amazon EC2. You might have databases that are running on SQL Server Enterprise edition but are a good fit for SQL Server Standard edition. You may also want to modernize your SQL Server database running on Windows to run on a Linux operating system to save on cost and licenses.

### Options

- [Microsoft SQL Server on Amazon EC2](#)
- [Amazon RDS for Microsoft SQL Server](#)
- [Amazon RDS Custom for SQL Server](#)

## Microsoft SQL Server on Amazon EC2

When to choose Microsoft SQL Server on Amazon EC2:

- You want full control over the database and access to its underlying operating system, database installation, and configuration.
- You want to administer your database, including backups and recovery, patching the operating system and the database, tuning the operating system and database parameters, managing security, and configuring high availability or replication.
- You want to use features and options that aren't currently supported by Amazon RDS. For more information, see [Features not supported and features with limited support](#) in the Amazon RDS documentation.
- You require a specific SQL Server version that isn't supported by Amazon RDS. For a list of supported versions and editions, see [SQL Server versions on Amazon RDS](#) in the *RDS for Microsoft SQL Server User Guide*.
- Your database size and performance requirements exceed the current RDS for Microsoft SQL Server offerings. For more information, see [Amazon RDS DB instance storage](#) in the *Amazon RDS User Guide*.
- You want to avoid automatic software patches that might not be compliant with your applications.
- You want to bring your own license instead of using the RDS for Microsoft SQL Server license-included model.
- You want to achieve higher IOPS and storage capacity than the current limits. For more information, see [Amazon RDS DB instance storage](#) in the *Amazon RDS User Guide*.

## Amazon RDS for Microsoft SQL Server

RDS for Microsoft SQL Server is a managed database service that simplifies the provisioning and management of SQL Server on AWS. With Amazon RDS, you can quickly deploy multiple versions and editions of SQL Server , with cost-efficient and resizeable compute capacity. You can provision

Amazon RDS for SQL Server DB instances with either General Purpose SSD or Provisioned IOPS SSD storage. Provisioned IOPS SSD is optimized for I/O-intensive, transactional (OLTP) database workloads.

Amazon RDS manages database administration tasks, including provisioning, backups, software patching, monitoring, and hardware scaling. Amazon RDS also offers Multi-AZ deployments and read replicas (for SQL Server Enterprise edition) to provide high availability, performance, scalability, and reliability for production workloads. For more information, see [Amazon RDS for Microsoft SQL Server](#).

When to choose RDS for Microsoft SQL Server:

- You want to focus on your business and applications, and you want AWS to take care of undifferentiated heavy lifting tasks, such as the provisioning of the database, management of backup and recovery tasks, management of security patches, minor SQL Server version upgrades, and storage management.
- You want a highly available database solution, and you want to take advantage of the push-button, synchronous Multi-AZ replication offered by Amazon RDS, without having to manually set up and maintain database mirroring, failover clusters, or Always On availability groups.
- You want to pay for the SQL Server license as part of the instance cost on an hourly basis, instead of making a large, up front investment.
- Your database size and IOPS requirements are supported by Amazon RDS for SQL Server. See [Amazon RDS DB Instance Storage](#) in the AWS documentation for the current maximum limits.
- You don't want to manage backups or point-in-time recoveries of your database.
- You want to focus on high-level tasks, such as performance tuning and schema optimization, instead of the daily administration of the database.
- You want to scale the instance type up or down based on your workload patterns without being concerned about licensing complexities.

## Amazon RDS Custom for SQL Server

Amazon RDS Custom for SQL Server is a managed database service for legacy, custom, and packaged applications that require access to the underlying operating system and database environment. Amazon RDS Custom for SQL Server automates setup, operation, and scaling of databases in the AWS Cloud while granting you access to the database and underlying operating system on Amazon EC2 to configure settings, install patches, and enable native features to meet

the dependent application's requirements. For more information, see [Working with RDS Custom for SQL Server](#) in the *Amazon Relational Database Service User Guide*.

When to choose Amazon RDS Custom for SQL Server:

- You want the benefits of Amazon RDS, but your requirements include the need to customize the underlying operating system and database environment for legacy, custom, and packaged applications.
- You need administrative rights to the database and underlying operating system.
- You need to install custom database and OS patches and packages.
- You need to configure file systems to share files directly with their applications.

## Microsoft SQL Server on Amazon EC2 concepts and terminology

The following concepts introduce you to the fundamental terminology used when working with Microsoft SQL Server on Amazon EC2 instances:

- [Amazon Machine Images \(AMIs\)](#)
- [Backup](#)
- [Billing](#)
- [High availability and disaster recovery \(HADR\)](#)
- [Instance](#)
- [Instance types](#)
- [Launching](#)
- [Security](#)
- [Storage](#)

### Amazon Machine Images (AMIs)

SQL Server on Amazon EC2 instances are created from Amazon Machine Images (AMIs). AMIs are similar to templates. SQL Server on Amazon EC2 AMIs are pre-installed with an operating system, typically Microsoft Windows Server, and other software. Together, these determine the operating

environment. You can select an AMI provided by AWS, create your own AMI, or select an AMI from the AWS Marketplace. To find a SQL Server on Amazon EC2 AMI, see the options under [Find a Windows AMI](#) in the *Amazon EC2 User Guide*.

## Backup

Your backup and recovery design for SQL Server on Amazon EC2 is flexible, depending on your RTO and RPO requirements. AWS provides the ability to perform server-level backups using Windows Volume Shadow Copy Service (VSS)-enabled Amazon Elastic Block Store (Amazon EBS) snapshots and with AWS Backup. You can also perform database-level backups using [native backup and restore procedures](#) for SQL Server databases. Database-level backups can be stored on Amazon EBS, FSx for Windows File Server, or Amazon Simple Storage Service using AWS Storage Gateway. For more information about backing up SQL Server on Amazon EC2, see [Backup and restore options for SQL Server on Amazon EC2](#) on the AWS Prescriptive Guidance website.

## Billing

A SQL Server on Amazon EC2 instance is charged by the second, with a minimum of 1 minute. Applied rates are based on the type and size of the selected instance, the edition of SQL Server when using a license-included instance, along with the cost of any additional services, such as storage or networking. AWS provides a variety of instance families that are favorable to the performance requirements of SQL Server workloads.

You can rent an instance based on your unique CPU, memory, and storage throughput requirements. You can also stop or terminate an instance at any time to pause or stop billing for the instance. The main advantage of the On-Demand model is the ability to save on CAPEX when an instance is no longer required.

### Warning

Any data on [Amazon EC2 instance store](#) volumes are lost if your instance is stopped or terminated. You'll still incur costs for EBS volumes when your instance is stopped. For more information, see [Stop and start your instance](#) in the *Amazon EC2 User Guide*.

## High availability and disaster recovery (HADR)

You can take advantage of Windows Server Failover Cluster for high availability and disaster recovery (HADR) with SQL Server on Amazon EC2. SQL Server on Amazon EC2 supports both

failover cluster instances (SQL FCIs) and Always On availability groups (AG). For more information see [How do I create a SQL Server Always On availability group cluster in the AWS Cloud?](#) in the AWS knowledge center.

## Instance

A SQL Server on Amazon EC2 instance is a virtual (or bare metal) server that runs in the AWS Cloud and can be provisioned on demand. The subscriber rents the virtual server by the hour/minute/second, and can use it to deploy specific configurations of SQL Server. For more information about On-Demand instances, see [On-Demand instances](#) in the *Amazon EC2 User Guide*.

An Amazon EC2 Dedicated Hosts is a physical server with EC2 instance capacity that is fully dedicated to your use. Dedicated Hosts allow you to use your existing per-socket, per-core, or per-VM Microsoft SQL Server software licenses. For more information about Dedicated Hosts, see [Dedicated Hosts](#) in the *Amazon EC2 User Guide*.

## Instance types

AWS provides various types of instances with different CPU, memory, storage, and networking configurations to support your application requirements. Each instance type is available in various sizes to address specific workload requirements. Instance types are grouped into families according to target application profiles, such as general purpose, compute-optimized, memory-optimized, and storage-optimized. The memory-optimized family of instances is a popular choice for SQL Server on Amazon EC2 because instances in this family have a high memory to CPU ratio for optimal performance. You can choose bare metal instances to support capabilities such as [Always Encrypted with secure enclaves on Amazon EC2 bare metal instances](#). For more information about individual and families of instance types, see [Amazon EC2 Instance Types](#) in the AWS product pages.

Some instance types support instance store volumes, which provide temporary block-level storage for the instance. If your instance has instance store volumes, we recommend that you place tempdb on an instance store volume. This can improve performance and decrease costs. For more information, see [Place tempdb in an instance store](#).

## Launching SQL Server on Amazon EC2

SQL Server on Amazon EC2 instances can be launched directly from the [Amazon EC2 console](#), with AWS CloudFormation, by using [AWS Tools for PowerShell](#), or by using the [AWS CLI](#). For a guided deployment of Microsoft SQL Server, use [AWS Launch Wizard](#).

## Security

AWS supports all security standards and compliance certifications, such as PCI-DSS, HIPAA/HITECH, FedRAMP, GDPR, FIPS, FIPS 140-2, and more. These standards enable you to build a fully compliant application on Amazon EC2. AWS also supports all SQL Server security features such as [Transparent Data Encryption](#) (TDE) and [Always Encrypted with Secure Enclaves](#) (when using bare metal instances).

Security and compliance is a shared responsibility between you and AWS. This shared model helps to relieve your operational burden because AWS operates, manages, and controls the components from the host operating system and virtualization layer to the physical security of the facilities in which the service operates.

For SQL Server on Amazon EC2, you assume responsibility and management of the guest operating system, including updates and security patches, other associated application software, and the configuration of AWS provided security group firewalls.

For more information about the shared responsibility model, see [Shared Responsibility Model](#).

## Storage

AWS provides many storage options to host your database files. In addition to EBS volume types, you can attach volumes to SQL Server on Amazon EC2 instances using an Amazon FSx managed file system service, such as FSx for Windows File Server and Amazon FSx for NetApp ONTAP. Some instance types provide an Amazon EC2 instance store which provides temporary block level storage on NVMe solid state drive (SSD) disks that are physically attached to the host computer. For more information, see [Best practices for deploying Microsoft SQL Server on Amazon EC2](#) on the AWS *Prescriptive Guidance* website.

## Best practices and recommendations for SQL Server clustering on Amazon EC2

You can configure Microsoft SQL Server on Amazon EC2 instances for high availability. SQL Server Always On availability groups offer high availability without the requirement for shared storage. The list of best practices in this topic, in addition to the prerequisites listed at [Prerequisites, Restrictions, and Recommendations for Always On availability groups](#), can help you optimize operating a SQL Server Always On availability groups on AWS. The practices listed in this topic also offer a method to gather logs.

**Note**

When nodes are deployed in different Availability Zones, or in different subnets within the same Availability Zone, they should be treated as a multi-subnet cluster. Keep this in mind as you apply these best practices and when you address possible failure scenarios.

**Contents**

- [Assign IP addresses](#)
- [Cluster properties](#)
- [Cluster quorum votes and 50/50 splits in a multi-site cluster](#)
- [DNS registration](#)
- [Elastic Network Adapters \(ENAs\)](#)
- [Multi-site clusters and EC2 instance placement](#)
- [Instance type selection](#)
- [Assign elastic network interfaces and IPs to the instance](#)
- [Heartbeat network](#)
- [Configure the network adapter in the OS](#)
- [IPv6](#)
- [Host record TTL for SQL Availability Group Listeners](#)
- [Logging](#)
- [NetBIOS over TCP](#)
- [NetFT Virtual Adapter](#)
- [Set possible owners](#)
- [Tune the failover thresholds](#)
- [Witness importance and Dynamic Quorum Architecture](#)
- [Troubleshoot](#)

**Assign IP addresses**

Each cluster node should have one elastic network interface assigned that includes three private IP addresses on the subnet: a primary IP address, a cluster IP address, and an Availability Group

IP address. The operating system (OS) should have the NIC configured for DHCP. It should not be set for a static IP address because the IP addresses for the cluster IP and Availability Group will be handled virtually in the Failover Cluster Manager. The NIC can be configured for a static IP as long as it is configured to only use the primary IP of **eth0**. If the other IPs are assigned to the NIC, it can cause network drops for the instance during failover events.

When the network drops because the IPs are incorrectly assigned, or when there is a failover event or network failure, it is not uncommon to see the following event log entries at the time of failure.

```
Isatap interface isatap.{9468661C-0AEB-41BD-BB8C-1F85981D5482} is no longer active.
```

```
Isatap interface isatap.{9468661C-0AEB-41BD-BB8C-1F85981D5482} with address  
fe80::5efe:169.254.1.105 has been brought up.
```

Because these messages seem to describe network issues, you could potentially mistake the cause of the outage or failure as a network error. However, these errors describe a symptom, rather than cause, of the failure. ISATAP is a tunneling technology that uses IPv6 over IPv4. When the IPv4 connection fails, the ISATAP adapter also fails. When the network issues are resolved, these entries should no longer appear in the event logs. Alternately, you can reduce network errors by safely disabling ISATAP with the following command.

```
netsh int ipv6 isatap set state disabled
```

When you run this command, the adapter is removed from Device Manager. This command should be run on all nodes. It does not impact the ability of the cluster to function. Instead, when the command has been run, ISATAP is no longer used. However, because this command might cause unknown impacts on other applications that use ISATAP, you should test it.

## Cluster properties

To see the complete cluster configuration, run the following PowerShell command.

```
Get-Cluster | Format-List -Property *
```

## Cluster quorum votes and 50/50 splits in a multi-site cluster

To learn how the cluster quorum works and what to expect if a failure occurs, see [Understanding Cluster and Pool Quorum](#).

## DNS registration

In Windows Server 2012, Failover Clustering, by default, attempts to register each DNS node under the cluster name. This is acceptable for applications that are aware the SQL target is configured for multi-site. However, when the client is not configured this way, it can result in timeouts, delays, and application errors due to attempts to connect to each individual node and failing on the inactive ones. To prevent these problems, the Cluster Resource parameter `RegisterAllProvidersIp` must be changed to **0**. For more information, see [RegisterAllProvidersIP Setting](#) and [Multi-subnet Clustered SQL + RegisterAllProvidersIP + SharePoint 2013](#).

The `RegisterAllProvidersIp` can be modified with the following PowerShell script.

```
Import-Module FailoverClusters
$cluster = (Get-ClusterResource | where {($_.ResourceType -eq "Network Name") -and
($_.OwnerGroup -ne "Cluster Group")}).Name
Get-ClusterResource $cluster | Set-ClusterParameter RegisterAllProvidersIP 0
Get-ClusterResource $cluster | Set-ClusterParameter HostRecordTTL 300
Stop-ClusterResource $cluster
Start-ClusterResource $cluster
```

In addition to setting the Cluster Resource parameter to **0**, you must ensure that the cluster has permissions to modify the DNS entry for your cluster name.

1. Log in to the Domain Controller (DC) for the domain, or a server that hosts the forward lookup zone for the domain.
2. Launch the DNS Management Console and locate the A record for the cluster.
3. Choose or right-click the A record, and choose **Properties**.
4. Choose **Security**.
5. Choose **Add**.
6. Choose **Object Types...**, select the box for **Computers**, and choose **OK**.
7. Enter the name of the cluster resource object and choose **Check name** and **OK if resolve**.
8. Select the check box for **Full Control**.
9. Choose **OK**.

## Elastic Network Adapters (ENAs)

AWS has identified known issues with some clustering workloads running on ENA driver version 1.2.3. We recommend upgrading to the latest version, and adjusting settings on the NIC in the operating system. For the latest versions, see [Amazon ENA Driver Versions](#). The first setting, which applies to all systems, increases Receive Buffers, which you can do with the following example PowerShell command.

```
Set-NetAdapterAdvancedProperty -Name (Get-NetAdapter | Where-Object {$_ .InterfaceDescription -like '*Elastic*'}).Name -DisplayName "Receive Buffers" -
DisplayValue 8192
```

For instances with more than 16 vCPUs, we recommend preventing RSS from running on CPU 0.

Run the following command.

```
Set-NetAdapterRss -name (Get-NetAdapter | Where-Object {$_ .InterfaceDescription -like '*Elastic*'}).Name -Baseprocessorgroup 0 -BaseProcessorNumber 1
```

## Multi-site clusters and EC2 instance placement

Each cluster is considered a [multi-site cluster](#). The EC2 service does not share IP addresses virtually. Each node must be in a unique [subnet](#). Though not required, we recommend that each node also be in a unique Availability Zone.

## Instance type selection

The type of instance recommended for Windows Server Failover Clustering depends on the workload. For production workloads, we recommend instances that support Amazon Elastic Block Store (Amazon EBS) optimization and enhanced networking. For more information, see [EBS optimization](#) and [Enhanced networking](#) in the *Amazon EC2 User Guide*.

## Assign elastic network interfaces and IPs to the instance

Each node in an EC2 cluster should have only one attached elastic network interface. The network interface should have a minimum of two assigned private IP addresses. However, for workloads that use Availability Groups, such as SQL Always On, you must include an additional IP address for each Availability Group. The primary IP address is used for accessing and managing the server, the

secondary IP address is used as the cluster IP address, and each additional IP address is assigned to Availability Groups, as needed.

## Heartbeat network

Some Microsoft documentation recommends using a dedicated [heartbeat network](#). However, this recommendation is not applicable to EC2. With EC2, while you can assign and use a second elastic network interface for the heartbeat network, it uses the same infrastructure and shares bandwidth with the primary network interface. Therefore, traffic within the infrastructure cannot be prioritized, and cannot benefit from a dedicated network interface.

## Configure the network adapter in the OS

The NIC in the OS can keep using DHCP as long as the DNS servers that are being retrieved from the DHCP Options Set allow for the nodes to resolve each other. You can set the NIC to be configured statically. When completed, you then manually configure only the primary IP address for the elastic network interface. Failover Clustering manages and assigns additional IP addresses, as needed.

For certain instance types, you can increase the maximum transmission unit (MTU) on the network adapter to support Jumbo Frames. This configuration reduces fragmentation of packets wherever Jumbo Frames are supported. For more information, see [Network maximum transmission unit \(MTU\) for your EC2 instance](#) in the *Amazon Elastic Compute Cloud User Guide*.

## IPv6

Microsoft does not recommend disabling IPv6 in a Windows Cluster. While Failover Clustering works in an IPv4-only environment, Microsoft tests clusters with IPv6 enabled. See [Failover Clustering and IPv6 in Windows Server 2012 R2](#) for details.

## Host record TTL for SQL Availability Group Listeners

Set the host record TTL to **300** seconds instead of the default 20 minutes (1200 seconds). For legacy client comparability, set `RegisterAllProvidersIP` to **0** for SQL Availability Group Listeners. This is not required in all environments. These settings are important because some legacy client applications cannot use `MultiSubnetFailover` in their connection strings. See [HostRecordTTL Setting](#) for more information. When you change these settings, the Cluster Resource must be restarted. The Cluster Group for the listener stops when the Cluster Resource

is restarted, so it must be started. If you do not start the Cluster Group, the Availability Group remains offline in a RESOLVING state. The following are example PowerShell scripts for changing the TTL and RegisterAllProvidersIP settings.

```
Get-ClusterResource yourListenerName | Set-ClusterParameter RegisterAllProvidersIP 0
```

```
Get-ClusterResource yourListenerName | Set-ClusterParameter HostRecordTTL 300
```

```
Stop-ClusterResource yourListenerName
```

```
Start-ClusterResource yourListenerName
```

```
Start-ClusterGroup yourListenerGroupName
```

## Logging

The default logging level for the cluster log is **3**. To increase the detail of log information, set the logging level to **5**. See [Set-ClusterLog](#) for more information about the PowerShell cmdlet.

```
Set-ClusterLog -Level 5
```

## NetBIOS over TCP

In Windows Server 2012 R2, you can increase the speed of the failover process by disabling NetBIOS over TCP. This feature was removed from Windows Server 2016. You should test this procedure if you are using earlier operating systems in your environment. For more information, see [Speeding Up Failover Tips-n-Tricks](#). The following is an example PowerShell command to disable NetBIOS over TCP.

```
Get-ClusterResource "Cluster IP Address" | Set-ClusterParameter EnableNetBIOS 0
```

## NetFT Virtual Adapter

For Windows Server versions earlier than 2016 and non-Hyper-V workloads, Microsoft recommends you enable the NetFT Virtual Adapter Performance Filter on the adapter in the OS.

When you enable the NetFT Virtual Adapter, internal cluster traffic is routed directly to the NetFT Virtual Adapter. For more information, see [NetFT Virtual Adapter Performance Filter](#). You can enable NetFT Virtual Adapter by selecting the check box in the NIC properties, or by using the following PowerShell command.


```
Get-NetAdapter | Set-NetAdapterBinding -ComponentID ms_netftflt -Enable $true
```

## Set possible owners

The Failover Cluster Manager can be configured so that each IP address specified on the Cluster Core Resources and Availability Group resources can be brought online only on the node to which the IP belongs. When the Failover Cluster Manager is not configured for this and a failure occurs, there will be some delay in failover as the cluster attempts to bring up the IPs on nodes that do not recognize the address. For more information, see [SQL Server Manages Preferred and Possible Owner Properties for AlwaysOn Availability Group/Role](#).

Each resource in a cluster has a setting for Possible Owners. This setting tells the cluster which nodes are permitted to “online” a resource. Each node is running on a unique subnet in a VPC. Because EC2 cannot share IPs between instances, the IP resources in the cluster can be brought online only by specific nodes. By default, each IP address that is added to the cluster as a resource has every node listed as a Possible Owner. This does not result in failures. However, during expected and unexpected failures, you can see errors in the logs about conflicting IPs and failures to bring IPs online. These errors can be ignored. If you set the Possible Owner property, you can eliminate these errors entirely, and also prevent down time while the services are moved to another node.

The following image shows an example of configuring an IP address so that it can only be brought online on the node to which the IP belongs:



## Tune the failover thresholds

In Windows Server 2012 R2, the network thresholds for the failover heartbeat network default to high values. See [Tuning Failover Cluster Network Thresholds](#) for details. This potentially unreliable configuration, which applies to clusters with some distance between them, was addressed in Server 2016 with an increase in the number of heartbeats. It was discovered that clusters would fail over due to very brief transient network issues. The heartbeat network is maintained with UDP 3343, which is traditionally far less reliable than TCP and more prone to incomplete conversations. Although there are low-latency connections between AWS Availability Zones, there are still geographic separations with a number of "hops" separating resources. Within an Availability Zone, there may be some distance between clusters unless the customer is using Placement Groups or Dedicated Hosts. As a result, there is a higher possibility for heartbeat failure with UDP than with TCP-based heartbeats.

The only time a cluster should fail over is when there is a legitimate outage, such as a service or node that experiences a hard failover, as opposed to a few UDP packets lost in transit. To ensure legitimate outages, we recommend that you adjust the thresholds to match, or even exceed, the settings for Server 2016 listed in [Tuning Failover Cluster Network Thresholds](#). You can change the settings with the following PowerShell commands.

```
(get-cluster).SameSubnetThreshold = 10
```

```
(get-cluster).CrossSubnetThreshold = 20
```

When you set these values, unexpected failovers should be dramatically reduced. You can fine tune these settings by increasing the delays between heartbeats. However, we recommend that you send the heartbeats more frequently with greater thresholds. Setting these thresholds even higher ensures that failovers occur only for hard failover scenarios, with longer delays before failing over. You must decide how much down time is acceptable for your applications.

After increasing the SameSubnetThreshold or CrossSubnetThreshold, we recommend that you increase the RouteHistoryLength to double the higher of the two values. This ensures that there is sufficient logging for troubleshooting. You can set the RouteHistoryLength with the following PowerShell command.

```
(Get-Cluster).RouteHistoryLength = 20
```

## Witness importance and Dynamic Quorum Architecture

There is a difference between Disk Witness and File Share Witness. Disk Witness keeps a backup of the cluster database while File Share Witness does not. Both add a [vote to the cluster](#). You can use Disk Witness if you use iSCSI-based storage. For more about witness options, see [File Share witness vs Disk witness for local clusters](#).

## Troubleshoot

If you experience unexpected failovers, first make sure that you are not experiencing networking, service, or infrastructure issues.

1. Check that your nodes are not experiencing network-related issues.

2. Check driver updates. If you are using outdated drivers on your instance, you should update them. Updating your drivers might address bugs and stability issues that might be present in your currently installed version.
3. Check for any possible resource bottlenecks that could cause an instance to become unresponsive, such as CPU and disk I/O. If the node cannot service requests, it might appear to be down by the cluster service.

# Set up Microsoft SQL Server on Amazon EC2

Describes the prerequisites, permissions, and configurations that you should consider when preparing to use Microsoft SQL Server on Amazon EC2 instances for your SQL Server workloads.

## Topics for setting up SQL Server on Amazon EC2

- [Prerequisites for using SQL Server on Amazon EC2](#)
- [Permissions required to use SQL Server on Amazon EC2](#)

## Prerequisites for using SQL Server on Amazon EC2

Complete the tasks in this section to start using SQL Server on Amazon EC2 instances for the first time:

1. [Sign up for an AWS account](#)
2. [Create a key pair](#)
3. [Create a security group](#)

## Sign up for an AWS account

To get started with AWS, you need an AWS account. For information about creating an AWS account, see [Getting started with an AWS account](#) in the *AWS Account Management Reference Guide*.

## Create a key pair

AWS uses public-key cryptography to secure the login information for your instance. You specify the name of the key pair when you launch your instance, then provide the private key to obtain the administrator password for your Windows instance so you can log in using RDP.

If you haven't created a key pair already, you can create one by using the Amazon EC2 console. Note that if you plan to launch instances in multiple Regions, you'll need to create a key pair in each Region. For more information about Regions, see [Regions and Zones](#) in the *User Guide for Windows Instances*.

## To create your key pair

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. In the navigation pane, choose **Key Pairs**.
3. Choose **Create key pair**.
4. For **Name**, enter a descriptive name for the key pair. Amazon EC2 associates the public key with the name that you specify as the key name. A key name can include up to 255 ASCII characters. It can't include leading or trailing spaces.
5. For **Key pair type**, choose either **RSA** or **ED25519**. Note that **ED25519** keys are not supported for Windows instances.
6. For **Private key file format**, choose the format in which to save the private key. To save the private key in a format that can be used with OpenSSH, choose **pem**. To save the private key in a format that can be used with PuTTY, choose **ppk**.

If you chose **ED25519** in the previous step, the **Private key file format** options do not appear, and the private key format defaults to **pem**.

7. Choose **Create key pair**.
8. The private key file is automatically downloaded by your browser. The base file name is the name you specified as the name of your key pair, and the file name extension is determined by the file format you chose. Save the private key file in a safe place.

### Important

This is the only chance for you to save the private key file.

For more information, see [Amazon EC2 key pairs and Windows instances](#) in the *User Guide for Windows Instances*.

## Create a security group

Security groups act as a firewall for associated instances, controlling both inbound and outbound traffic at the instance level. You must add rules to a security group that enable you to connect to your instance from your IP address using RDP. You can also add rules that allow inbound and outbound HTTP and HTTPS access from anywhere.

Note that if you plan to launch instances in multiple Regions, you'll need to create a security group in each Region. For more information about Regions, see [Regions and Zones](#) in the *User Guide for Windows Instances*.

## Prerequisites

You'll need the public IPv4 address of your local computer. The security group editor in the Amazon EC2 console can automatically detect the public IPv4 address for you. Alternatively, you can use the search phrase "what is my IP address" in an Internet browser, or use the following service: [Check IP](#). If you are connecting through an Internet service provider (ISP) or from behind a firewall without a static IP address, you need to find out the range of IP addresses used by client computers.

You can create a custom security group using one of the following methods.

### New Amazon EC2 console

#### To create a security group with least privilege

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. From the top navigation bar, select a Region for the security group. Security groups are specific to a Region, so you should select the same Region in which you created your key pair.
3. In the left navigation pane, choose **Security Groups**.
4. Choose **Create security group**.
5. For **Basic details**, do the following:
  - a. Enter a name for the new security group and a description. Use a name that is easy for you to remember, such as your user name, followed by `_SG_`, plus the Region name. For example, `me_SG_uswest2`.
  - b. In the **VPC** list, select your default VPC for the Region.
6. For **Inbound rules**, create rules that allow specific traffic to reach your instance. For example, use the following rules for a web server that accepts HTTP and HTTPS traffic. For more examples, see [Security group rules for different use cases](#) in the *User Guide for Windows Instances*.
  - a. Choose **Add rule**. For **Type**, choose **HTTP**. For **Source**, choose **Anywhere**.
  - b. Choose **Add rule**. For **Type**, choose **HTTPS**. For **Source**, choose **Anywhere**.
  - c. Choose **Add rule**. For **Type**, choose **RDP**. For **Source**, do one of the following:

- Choose **My IP** to automatically add the public IPv4 address of your local computer.
- Choose **Custom** and specify the public IPv4 address of your computer or network in CIDR notation. To specify an individual IP address in CIDR notation, add the routing suffix /32, for example, 203.0.113.25/32. If your company or your router allocates addresses from a range, specify the entire range, such as 203.0.113.0/24.

 **Warning**

For security reasons, do not choose **Anywhere** for **Source** with a rule for RDP. This would allow access to your instance from all IP addresses on the internet. This is acceptable for a short time in a test environment, but it is unsafe for production environments.

7. For **Outbound rules**, keep the default rule, which allows all outbound traffic.
8. Choose **Create security group**.

## Old Amazon EC2 console

### To create a security group with least privilege

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. In the left navigation pane, choose **Security Groups**.
3. Choose **Create Security Group**.
4. Enter a name for the new security group and a description. Use a name that is easy for you to remember, such as your user name, followed by \_SG\_, plus the Region name. For example, *me\_SG\_uswest2*.
5. In the **VPC** list, select your default VPC for the Region.
6. On the **Inbound rules** tab, create the following rules (choose **Add rule** for each new rule):
  - Choose **HTTP** from the **Type** list, and make sure that **Source** is set to **Anywhere** (0.0.0.0/0).
  - Choose **HTTPS** from the **Type** list, and make sure that **Source** is set to **Anywhere** (0.0.0.0/0).

- Choose **RDP** from the **Type** list. In the **Source** box, choose **My IP** to automatically populate the field with the public IPv4 address of your local computer. Alternatively, choose **Custom** and specify the public IPv4 address of your computer or network in CIDR notation. To specify an individual IP address in CIDR notation, add the routing suffix /32, for example, 203.0.113.25/32. If your company allocates addresses from a range, specify the entire range, such as 203.0.113.0/24.

 **Warning**

For security reasons, do not allow RDP access from all IP addresses to your instance. This is acceptable for a short time in a test environment, but it is unsafe for production environments.

7. On the **Outbound rules** tab, keep the default rule, which allows all outbound traffic.
8. Choose **Create security group**.

## Command line

### To create a security group with least privilege

Use one of the following commands:

- [create-security-group](#) (AWS CLI)
- [New-EC2SecurityGroup](#) (AWS Tools for Windows PowerShell)

For more information, see [Amazon EC2 security groups for Windows instances](#) in the *Amazon EC2 User Guide*.

## Permissions required to use SQL Server on Amazon EC2

For information about the permissions required to create or modify Amazon EC2 resources, or to perform tasks using the Amazon EC2 API, see [IAM policies for Amazon EC2](#) in the *User Guide for Windows Instances*.

# Understand licensing options and considerations for Microsoft SQL Server on Amazon EC2

There are two ways in which you can license Microsoft SQL Server on Amazon EC2 on the AWS Cloud. You acquire your own existing SQL Server licenses, or those which are provided by AWS. The most cost-effective license strategy for your workload will depend on multiple factors. For more information on comparing the costs of SQL Server editions, see [Compare SQL Server editions](#) on the AWS Prescriptive Guidance website.

## Topics

- [Licensing options](#)
- [Licensing considerations](#)
- [Amazon EC2 High Availability for SQL Server on Amazon EC2](#)

## Licensing options

You can launch Amazon Elastic Compute Cloud (Amazon EC2) instances with Microsoft SQL Server licenses included from AWS, or you can bring your own SQL Server licenses for use on AWS. You can perform a license type conversion for SQL Server in certain configurations if your needs change. For the most license flexibility, you can import your VM into AWS. For more information, see [Eligible license types for license type conversion](#) in the *AWS License Manager User Guide*.

### Licensing options topics

- [License-included](#)
- [BYOL](#)

## License-included

Windows Server with currently supported versions of Microsoft SQL Server AMIs are available from AWS in a variety of combinations. AWS provides these AMIs with SQL Server software and operating system updates already installed. When you purchase an Amazon EC2 instance with a Windows Server AMI, licensing costs and compliance are handled for you. For more information, see [Find a SQL Server license-included AMI](#).

Amazon EC2 offers a variety of instance types and sizes that you can configure for your target workload. Amazon EC2 AMIs with Windows Server require no Client Access Licenses (CALs). They also include two Microsoft Remote Desktop Services licenses for administrative purposes.

For SQL Server license-included AMIs, use the installation and setup media included in C:\SQLServerSetup to perform in-place SQL Server version upgrades, make changes to the default installation, add new features, or install additional named instances.

## BYOL

When you launch a SQL Server instance from an imported AMI, you can bring your existing licenses with the Bring Your Own License model (BYOL), and let AWS manage them to ensure compliance with licensing rules that you set. To import your own licensed image, you can use a service such as [VM Import/Export](#) or [AWS Transform MGN](#). After you import your licensed image, and it is available as a private AMI in your AWS account on the Amazon EC2 console, you can use the AWS License Manager service to create a license configuration.

After you create the license configuration, you must associate the AMI that contains your licensed operating system image with the configuration. Then, you must create a host resource group and associate it with the license configuration. After you associate your host resource group with the configuration, License Manager automatically manages your hosts when you launch instances into a host resource group, and ensures that you do not exceed your configured license count limits. For more information, see the [Getting started](#) section of the *License Manager User Guide*.

You can also bring your own SQL Server licenses with Active Software Assurance to default (shared) tenant Amazon EC2 through Microsoft License Mobility through Software Assurance. For information about how to sign up for Microsoft License Mobility, see [License Mobility](#).

## Licensing considerations

There are many considerations for cost effectively licensing your Microsoft SQL Server on Amazon EC2 workload. Your use case, and existing license agreements, will determine whether to bring your own license to AWS with the Bring Your Own License model (BYOL) or to use license included AMIs from AWS. The following topics should help determine which approach you might take. For more information, see [Licensing - SQL Server](#) on the *Amazon Web Services and Microsoft Frequently Asked Questions* page.

### Licensing considerations topics

- [Choose a SQL Server edition](#)

- [Purchase SQL Server from AWS](#)
- [Use BYOL for SQL Server on AWS](#)
- [Quantify the required SQL Server licenses for BYOL](#)
- [License Mobility with SQL Server](#)
- [Track BYOL license consumption](#)
- [SQL Server client access licenses \(CALs\)](#)
- [Licensing for passive failover](#)

## Choose a SQL Server edition

The edition of SQL Server that is used will determine the supported features your implementation will have available. For example, the edition determines the maximum compute capacity used by a single instance of the SQL Server Database Engine, and the high availability options you might implement. For a comparison of SQL Server editions and supported features, see [Editions and supported features of SQL Server 2022](#) in the Microsoft documentation.

## Purchase SQL Server from AWS

You can utilize Microsoft SQL Server licenses included from AWS. You can choose any of the following editions for your use on Amazon EC2 instances.

- SQL Server Web
- SQL Server Standard
- SQL Server Enterprise

### Note

- SQL Server Express AMIs are available for use from AWS. This free edition of SQL Server doesn't incur additional charges as there is no licensing fee.
- SQL Server Developer edition is eligible for use in non-production, development, and test workloads. Once downloaded from Microsoft, you can bring and install SQL Server Developer edition on Amazon EC2 instances in the AWS Cloud. Dedicated infrastructure is not required for SQL Server Developer edition. For more information, see <https://www.microsoft.com/en-us/sql-server/sql-server-downloads>.

## Use BYOL for SQL Server on AWS

You can use BYOL licenses for SQL Server on AWS. The requirements differ depending on if the licenses have active Software Assurance.

### SQL Server licenses with active Software Assurance

You can bring your SQL Server licenses with active Software Assurance to default (shared) tenant Amazon EC2 through License Mobility benefits. Microsoft requires that you complete and send a License Mobility verification form which can be downloaded [here](#). For more information, see [License Mobility](#).

### SQL Server licenses without active Software Assurance

SQL Server licenses without Software Assurance can be deployed on Amazon Elastic Compute Cloud Dedicated Hosts if the licenses are purchased prior to 10/1/2019 or added as a true-up under an active Enterprise Enrollment that was effective prior to 10/1/2019. In these specific BYOL scenarios, the licenses can only be upgraded to versions that were available prior to 10/1/2019. For more information, see [Dedicated Hosts](#) in the *Amazon EC2 User Guide*, and the [Amazon EC2 Dedicated Hosts FAQs](#).

## Quantify the required SQL Server licenses for BYOL

If you are licensing SQL Server under Microsoft License Mobility through Software Assurance, the number of licenses required varies based on the instance type, version of SQL Server, and the Microsoft licensing model you choose. For assistance with virtual core licensing calculations under the Microsoft Product Terms based on the instance type, see [SQL License Mobility](#).

If you are using Dedicated Hosts, Amazon EC2 provides you with the number of physical cores installed on the Dedicated Host. Using this information, you can calculate the number of SQL Server licenses that you need to bring in. For more information, see [Amazon EC2 Dedicated Hosts Pricing](#) and the [SQL Server 2022 licensing guide](#).

## License Mobility with SQL Server

SQL Server licenses with active Software Assurance are eligible for Microsoft License Mobility and can be deployed on default or dedicated tenant Amazon EC2. For more information on bringing SQL Server licenses with active Software Assurance to default tenant EC2, see [Microsoft License Mobility](#).

It is also possible to bring SQL Server licenses without active Software Assurance to EC2 Dedicated Hosts. To be eligible, the licenses must be purchased prior to October 1, 2019 or added as a true-up under an active Enterprise Enrollment that was effective prior to October 1, 2019. For additional FAQs about Dedicated Hosts, see the [Dedicated Hosts](#) section of the *Amazon Web Services and Microsoft FAQ*.

## Track BYOL license consumption

You can use AWS License Manager to manage your software licenses for SQL Server. With License Manager, you can create license configurations, take inventory of your license-consuming resources, associate licenses with resources, and track inventory and compliance. For more information, see [What is AWS License Manager?](#) in the *AWS License Manager User Guide*.

## SQL Server client access licenses (CALs)

When you are using SQL Server on Amazon EC2, license included instances do not require client access licenses (CALs) for SQL Server. An unlimited number of end users can access SQL Server on a license-included instance.

When you bring your own SQL Server licenses to Amazon EC2 through Microsoft License Mobility or BYOL, you must continue to follow the licensing rules in place on-premises. If you purchased SQL Server under the Server/CAL model, you still require CALs to meet Microsoft licensing requirements, but these CALs would remain on-premises and enable end user access SQL Server running on AWS.

## Licensing for passive failover

There are various factors to consider when licensing passive failover for SQL Server. The information in this section pertains only to the SQL Server licenses and not the Windows Server licenses. In all cases, you must license Windows Server.

### Using instances that include the license for SQL Server

When you purchase SQL Server license included instances on EC2, you must license passive failover instances.

### Bringing SQL Server licenses with active Software Assurance to default tenant Amazon EC2

When you bring SQL Server 2014 and later versions with Software Assurance to default tenant EC2, you must license the virtual cores (vCPUs) on the active instance. In return, Software

Assurance permits one passive instance (equal or lesser size) where SQL Server licensing is not required.

## Bringing SQL Server to Amazon EC2 Dedicated Instances

SQL Server 2014 and later versions require Software Assurance for SQL Server passive failover benefits on dedicated infrastructure. When you bring SQL Server with Software Assurance, you must license the cores on the active instance/host and are permitted one passive instance/host (equal or lesser size) where SQL Server licensing is not required.

SQL Server 2008 - SQL Server 2012R2 are eligible for passive failover on an Amazon EC2 Dedicated Hosts infrastructure without active Software Assurance. In these scenarios, you will license the active instance/host, and it will be permitted one passive instance/host of equal or lesser size where SQL Server licensing is not required.

There are specific BYOL scenarios that do not require Microsoft License Mobility through Software Assurance. An Amazon EC2 Dedicated Hosts infrastructure is always required in these scenarios. To be eligible, the licenses must be purchased prior to October 1, 2019 or added as a true-up under an active Enterprise Enrollment that was effective prior to October 1, 2019. In these specific BYOL scenarios, the licenses can only be upgraded to versions that were available prior to October 1, 2019.

## Amazon EC2 High Availability for SQL Server on Amazon EC2

Amazon EC2 High Availability for SQL Server (SQL HA) allows you to configure Amazon EC2 instances running license-included SQL Server as part of a HA cluster to reduce licensing costs. This feature identifies the SQL Server standby (also known as passive) instances in your SQL HA deployments and waives SQL Server licensing fees for them, allowing you to pay only Windows rates for these standby instances while maintaining your HA configuration.

### Supported SQL Server High Availability deployments

SQL HA supports two SQL Server High Availability deployments, including:

- Always On Availability Groups
- Always On Failover Cluster Instances

For Always On Availability Groups, SQL HA identifies primary and secondary (also known as standby or passive) replicas and waives the SQL Server licenses on the secondary replicas (Amazon

EC2 instances) that are not actively serving read traffic. For Failover Cluster Instances, SQL HA recognizes which instances are active and which are standing by for failover scenarios and waives the SQL Server licenses on the standby instances. For more information about the SQL HA configurations, see [Deploy SQL Server on Amazon EC2](#)

## Considerations and requirements

- SQL HA supports two Amazon EC2 instances (also known as nodes) per SQL Server HA cluster.
- The SQL HA active instance should have equal or more vCPUs than the standby instance.
- SQL HA saves license costs for SQL Server license-included only. For more information, see [SQL Server licensing options](#).
- Cross-Region deployments are not supported.
- The SQL Server standby node must meet requirements to receive the license savings, including:
  1. Does not serve incoming traffic;
  2. Does not run active SQL Server workloads;
  3. Is not a readable secondary (except master, msdb, tempdb, or model databases);
  4. For Always On Availability Groups, there is no standalone database running outside of the Availability Group.
- SQL HA supports SQL Server (Standard and Enterprise Editions) 2017 and later on Windows Server 2019 and later.
- Windows PowerShell must be 5.1 or above.
- Amazon EC2 Reserved Instances are not supported by SQL HA. If you are using Reserved Instances, discounts may not be applied to these instances in the same payer account. We recommend using Savings Plans instead to benefit from both SQL HA license savings and Savings Plans compute cost savings. For more information, see the [Savings Plans User Guide](#).
- This feature may be terminated at any time, in which case AWS will provide you with as much prior notice as is reasonably practicable under the circumstances.

## Prerequisites

To get the license savings for your SQL HA workload, your environment must meet several requirements:

- You have SQL Server license-included HA workloads on Amazon EC2. You can use AWS Launch Wizard to simplify the SQL Server deployment on Amazon EC2. For more information, see [AWS Launch Wizard for SQL Server](#).

- AWS Systems Manager Agent (SSM Agent) must be installed and running on the instances in the SQL HA deployment. For more information, see [Working with SSM Agent](#).
- You must attach the **AWSEC2SqlHaInstancePolicy** managed policy to your instance role or use a custom role with the required permissions for the Amazon EC2 to access your instances, including permissions for AWS Systems Manager to run commands, access to AWS Secrets Manager for retrieving SQL Server credentials, and Amazon EC2 permissions to read instance metadata.
- By default, AWS Systems Manager uses the built-in [NT AUTHORITY\SYSTEM] user to access SQL Server HA metadata. You will need to store secrets in AWS Secrets Manager only when your security policies have restricted or disabled the [NT AUTHORITY\SYSTEM] account.
- Network connectivity allows AWS Systems Manager commands to reach your instances.

## How Amazon EC2 High Availability for SQL Server works

Upon registration, Amazon EC2 High Availability for SQL Server (SQL HA) automatically monitors your Amazon EC2 instances running Windows SQL Server License Included AMIs and classifies them as either active or standby based on their current role in your SQL Server deployment. For High Availability configurations containing an active SQL Server instance, one standby failover instance in the same cluster may receive a SQL Server licensing fee waiver, meaning you pay only the Windows Server licensing fee. You can monitor your current SQL HA status through the Amazon EC2 console, which displays the latest records of which instances are receiving license savings and historical status changes.

SQL HA continuously monitors your enabled SQL Server instances to determine their active or standby status. Using AWS Systems Manager (SSM) commands, it collects metadata from your SQL Server installations and applies classification logic to identify which instances are actively serving traffic and which are functioning as standby failover nodes.

Standby instances are billed as Windows instances rather than Windows SQL Server instances, providing license cost savings. Billing changes take effect when an SQL HA standby detection enabled instance is classified as standby and eligible for the benefit, with no manual intervention required. This classification adapts to changes in your environment, such as failover events where a standby instance becomes active. The system detects these transitions and updates billing accordingly.

# Getting started with Amazon EC2 High Availability for SQL Server

To get started with Amazon EC2 High Availability for SQL Server (SQL HA), perform the following steps:

## Topics

- [Step 1: Set up SSM Agent](#)
- [Step 2: Attach AWS managed policy to instances](#)
- [Step 3: \(Optional\) Store SQL Server credentials in AWS Secrets Manager](#)
- [Step 4: EnableSQL HA license savings](#)
- [Windows user setup for Amazon EC2 High Availability for SQL Server](#)

## Step 1: Set up SSM Agent

The Systems Manager Agent (SSM Agent) must be installed and running on the Amazon EC2 SQL Server instances with the High Availability deployments. The SSM Agent executes an SSM document to determine and report the SQL HA state for the instance.

The SSM Agent is preinstalled, by default, on the Amazon Machine Images (AMIs) for Windows and SQL Server provided by Amazon. For more information, see [AWS Windows AMIs](#). To check if SSM Agent is correctly configured on your instances, you can use the System Manager console, or call [DescribeInstanceInformation](#) to verify the SSM Agent [PingStatus](#) is Online. If necessary, you can manually download and install the latest version of SSM Agent on your Amazon EC2 SQL Server instances. For more information, see [Manually install the SSM Agent on Amazon EC2 instances for Windows Server](#).

## Step 2: Attach AWS managed policy to instances

To ensure that your instance has the required IAM permissions, you must attach the following AWS managed policies to the instance:

- **AWSEC2SqlHaInstancePolicy** — grants permissions for SQL HA to execute AWS Systems Manager (SSM) Run Command document AWSEC2-DetectSqlHaState to automatically detect the standby state of your SQL Server instances.
- **AmazonSSMManagedInstanceCore** — enables AWS Systems Manager service core functionality.

For more information, see [Attach an IAM role to an Amazon EC2 instance](#).

**Note**

If needed, you can create and attach your own custom IAM role. However, at a minimum, the role must include all of the permissions that are included in the **AWSEC2SqlHaInstancePolicy** AWS managed policy.

### Step 3: (Optional) Store SQL Server credentials in AWS Secrets Manager

**By default**, AWS Systems Manager uses the built-in [NT AUTHORITY\SYSTEM] user to access SQL Server HA metadata. If you choose to use the built-in [NT AUTHORITY\SYSTEM] user, you may need to configure Windows user permissions to ensure the service can obtain High Availability metadata from your SQL Server instances. For more information, see [Windows user setup for Amazon EC2 High Availability for SQL Server](#).

**Alternatively**, if your security policies have restricted or disabled the [NT AUTHORITY\SYSTEM] account, you will need to store and use your SQL Server credentials in AWS Secrets Manager. For more information, see [Create a secret in AWS Secrets Manager with appropriate SQL Server permissions](#).

### Step 4: Enable SQL HA license savings

You must enable SQL HA standby detection for Windows SQL Server license-included instances to receive SQL Server license savings. Use one of the following methods:

#### Console

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. In the navigation panel, choose **Instances**.
3. Select the instances in the High Availability deployment to enable SQL HA standby detection monitoring, choose **Actions**, **Instance settings**, **Modify SQL High Availability settings**.
4. In the **Review prerequisites** step, review each instance to make sure it is configured correctly.
  - The **SSM agent status** column indicates the state of the SSM Agent on the instance. **Online** indicates that the SSM Agent is running and accessible.

- The **Recommended IAM policies** column indicates whether the instance has an attached IAM role with the required permissions. We recommend attaching the service managed policy `AWSEC2SqlHaInstancePolicy` to the instance or you can use any equivalent custom inline policy. **Verified** indicates that the instance has the managed policy attached while it doesn't verify the permission if you use other custom policies. The **IAM role** column indicates the currently attached IAM role. To attach a different role, choose **Modify IAM role**.
5. Choose **Next**.
  6. In the **Manage SQL High Availability license savings** step, for each instance do the following:
    - For **SQL High Availability license savings**, select **Enable**.
    - (Optional) For **SQL Server credentials**, select the secret that has the SQL Server credentials for that instance .
  7. Choose **Next**.
  8. In the **Review and apply changes** step, review the configuration and then choose **Apply changes**.

## AWS CLI

Use the [enable-instance-sql-ha-standby-detections](#) command. For `instance-ids` specify the IDs of the instances to opt in. If you choose to perform Step 3: Create secret for SQL Server credentials, specify the optional `--sql-server-credentials` with the Amazon Web Services secret arn that has the SQL Server credentials in.

```
aws ec2 enable-instance-sql-ha-standby-detections \  
--instance-ids instance_ids \  
--sql-server-credentials secret_manager_secret_arn
```

## Windows user setup for Amazon EC2 High Availability for SQL Server

### Note

You only need to perform the steps in this section if you choose to use the default, built-in `[NT AUTHORITY\SYSTEM]` user as described in [Step 3: Store SQL Server credentials in](#)

[AWS Secrets Manager](#). If you choose to store custom SQL Server credentials in AWS Secrets Manager, these Windows user setup steps are not required.

Amazon EC2 High Availability for SQL Server uses AWS Systems Manager (SSM) to connect to Amazon EC2 instances and obtain SQL Server High Availability metadata. The SSM command runs under the context of the default local user on the Amazon EC2 instance: NT AUTHORITY\SYSTEM. If you performed post-launch lockdowns on your SQL Server instances by removing certain default SQL Server permissions and built-in groups, you may need to perform a few steps to grant required permissions to NT AUTHORITY\SYSTEM.

Additionally, when enabling your Amazon EC2 instances for SQL HA standby detection, you can optionally provide an AWS secret containing credentials to a Windows domain user or local user on your Amazon EC2 instances other than the default local user, NT AUTHORITY\SYSTEM. The service uses this provided Windows user to connect to all SQL Server instances on the Amazon EC2 instance and run SQL Server queries to obtain High Availability metadata. This guide explains how to either grant required permissions to NT AUTHORITY\SYSTEM, or how to create a Windows domain or local user with least permissions required for the service to process Amazon EC2 instances enabled for SQL HA standby detection, and how to create an AWS secret containing credentials for this user.

## Topics

- [Option 1: Grant required permissions to the \[NT AUTHORITY\SYSTEM\] user](#)
- [Option 2: Create new domain user with required permissions](#)
- [Option 3: Create new local user with require permissions](#)

### Option 1: Grant required permissions to the [NT AUTHORITY\SYSTEM] user

This section covers the most straightforward setup to begin enabling Amazon EC2 instances for SQL HA standby detection. If you follow this section, you need not provide an AWS secret when enabling your Amazon EC2 instances for SQL HA standby detection, since SSM will use the default local user NT AUTHORITY\SYSTEM to authenticate into SQL Server.

Since SQL Server license-included AMIs allow NT AUTHORITY\SYSTEM to authenticate into SQL Server by default, the following steps may not be required to enable your instances for SQL HA standby detection. However, if you performed post-launch lockdowns on your SQL Server

instances, you may need to grant least permissions back to NT AUTHORITY\SYSTEM for the service to obtain High Availability metadata.

## To grant SQL Server access for NT AUTHORITY\SYSTEM

The following steps need to be repeated on every SQL Server instance on the Amazon EC2 instance. Amazon EC2 SQL HA obtains High Availability metadata on all SQL Server installs on the Amazon EC2 instance, so the default local user needs to be able to query SQL Server across all SQL Server instances.

- Connect to your Amazon EC2 instance and open SQL Server Management Studio, then run the following SQL Server command on each SQL Server instance. This command creates a SQL Server login for NT AUTHORITY\SYSTEM and grants minimal read-only SQL Server permissions for this user.

```
-- Create SQL Server login for default local user
IF NOT EXISTS (SELECT name FROM master.sys.server_principals WHERE name = 'NT
  AUTHORITY\SYSTEM')
BEGIN
    CREATE LOGIN [NT AUTHORITY\SYSTEM] FROM WINDOWS WITH DEFAULT_DATABASE=[master]
END

USE [master]
GO

IF NOT EXISTS (SELECT name FROM master.sys.sysusers WHERE name = 'NT AUTHORITY
  \SYSTEM')
BEGIN
    CREATE USER [NT AUTHORITY\SYSTEM] FOR LOGIN [NT AUTHORITY\SYSTEM]
END
GO

-- Grant database permissions
USE [master]
GO

IF NOT EXISTS (SELECT name FROM master.sys.database_principals WHERE name =
  'db_role_ec2_sql_ha')
BEGIN
    CREATE ROLE [db_role_ec2_sql_ha]
END
```

```
GRANT VIEW DATABASE STATE to [db_role_ec2_sql_ha]
GO

ALTER ROLE [db_role_ec2_sql_ha] ADD MEMBER [NT AUTHORITY\SYSTEM]
GO

-- Grant server permissions
USE [master]
GO

IF NOT EXISTS (SELECT name FROM master.sys.server_principals WHERE name =
    'svr_role_ec2_sql_ha')
BEGIN
    CREATE SERVER ROLE [svr_role_ec2_sql_ha]
END

GRANT VIEW SERVER STATE TO [svr_role_ec2_sql_ha]
GRANT VIEW ANY DEFINITION TO [svr_role_ec2_sql_ha]
GRANT VIEW ANY DATABASE TO [svr_role_ec2_sql_ha]
GO

ALTER SERVER ROLE [svr_role_ec2_sql_ha] ADD MEMBER [NT AUTHORITY\SYSTEM]
GO
```

Your default local user setup is complete. You can now enable SQL HA standby detection for your Amazon EC2 instances.

## Option 2: Create new domain user with required permissions

This section covers how to create a Windows domain user with the necessary permissions to connect to SQL Server and obtain High Availability metadata. This option is preferred over creating a new local user, as the domain user can be used for any Amazon EC2 instance joined to the domain. This allows you to supply just one AWS secret for multiple Amazon EC2 instances enabled for SQL HA standby detection.

### To create and configure a domain user

#### 1. Create a domain user

This step differs based on the type of Active Directory (AD) being used, and assumes the Amazon EC2 instances you wish to enable for SQL HA standby detection are already joined to

this domain. For an AWS managed Microsoft AD, use the following AWS CLI commands to create a new domain user. Replace *username* and *password* with your desired username and password.

```
aws ds-data create-user \  
  --directory-id directory-id \  
  --sam-account-name "username"
```

Then assign a password to the domain user:

```
aws ds reset-user-password \  
  --directory-id directory-id \  
  --user-name "username" \  
  --new-password "password"
```

## 2. Create an AWS secret containing credentials

Save the Windows domain user credentials to an AWS secret. The domain user's username must be saved in the following format: *directory-netBIOS-name\username*. The *directory-netBIOS-name* is the directory NetBIOS name of your AD.

```
aws secretsmanager create-secret \  
  --name "domain-user-credentials" \  
  --description "Domain user credentials for EC2 SQL HA standby detection." \  
  --secret-string "{\"username\": \"directory-netBIOS-name\\username\", \"password\": \"password\"}"
```

## 3. Grant SQL Server access for domain user

Connect to your Amazon EC2 instance and open SQL Server Management Studio, then run the following SQL Server command on each SQL Server instance. Replace *username* with the username you selected and *directory-netBIOS-name* with the AD's directory NetBIOS name.

```
-- Create SQL Server login for domain user  
IF NOT EXISTS (SELECT name FROM master.sys.server_principals WHERE name =  
  'directory-netBIOS-name\\username')  
BEGIN  
  CREATE LOGIN [directory-netBIOS-name\\username] FROM WINDOWS WITH  
  DEFAULT_DATABASE=[master]  
END
```

```
USE [master]
GO

IF NOT EXISTS (SELECT name FROM master.sys.sysusers WHERE name = 'directory-
netBIOS-name\username')
BEGIN
    CREATE USER [directory-netBIOS-name\username] FOR LOGIN [directory-netBIOS-
name\username]
END
GO

-- Grant database permissions
USE [master]
GO

IF NOT EXISTS (SELECT name FROM master.sys.database_principals WHERE name =
'db_role_ec2_sql_ha')
BEGIN
    CREATE ROLE [db_role_ec2_sql_ha]
END

GRANT VIEW DATABASE STATE to [db_role_ec2_sql_ha]
GO

ALTER ROLE [db_role_ec2_sql_ha] ADD MEMBER [directory-netBIOS-name\username]
GO

-- Grant server permissions
USE [master]
GO

IF NOT EXISTS (SELECT name FROM master.sys.server_principals WHERE name =
'svr_role_ec2_sql_ha')
BEGIN
    CREATE SERVER ROLE [svr_role_ec2_sql_ha]
END

GRANT VIEW SERVER STATE TO [svr_role_ec2_sql_ha]
GRANT VIEW ANY DEFINITION TO [svr_role_ec2_sql_ha]
GRANT VIEW ANY DATABASE TO [svr_role_ec2_sql_ha]
GO
```

```
ALTER SERVER ROLE [svr_role_ec2_sql_ha] ADD MEMBER [directory-netBIOS-  
name\username]  
GO
```

Your domain user setup is complete. When enabling Amazon EC2 instances for SQL HA standby detection, you can supply the ARN for the AWS secret you created.

### Option 3: Create new local user with require permissions

This section covers how to create a Windows local user restricted to a single Amazon EC2 instance with the necessary permissions to connect to SQL Server and obtain High Availability metadata.

#### To create and configure a local user

##### 1. Create a local user on the Amazon EC2 instance

Connect to your Amazon EC2 instance and open PowerShell as Administrator, then execute the following command. Replace *username* and *password* with your desired username and password.

```
New-LocalUser -Name "username" -Password (ConvertTo-SecureString "password" -  
AsPlainText -Force) -Description "Local user for EC2 SQL HA standby detection."
```

##### 2. Create an AWS secret containing credentials

Save the Windows local user credentials to an AWS secret.

```
aws secretsmanager create-secret \  
  --name "local-user-credentials" \  
  --description "Local user credentials for EC2 SQL HA standby detection." \  
  --secret-string "{\"username\":\"username\",\"password\":\"password\"}"
```

##### 3. Grant SQL Server access for local user

Connect to your Amazon EC2 instance and open SQL Server Management Studio, then run the following SQL Server command on each SQL Server instance. Replace *username* with the username you selected and *COMPUTERNAME* with the Amazon EC2 instance computer name. You can retrieve the computer name with the PowerShell command `$env:COMPUTERNAME`.

```
-- Create SQL Server login for local user
```

```
IF NOT EXISTS (SELECT name FROM master.sys.server_principals WHERE name =
'COMPUTERNAME\username')
BEGIN
    CREATE LOGIN [COMPUTERNAME\username] FROM WINDOWS WITH
    DEFAULT_DATABASE=[master]
END

USE [master]
GO

IF NOT EXISTS (SELECT name FROM master.sys.sysusers WHERE name =
'COMPUTERNAME\username')
BEGIN
    CREATE USER [COMPUTERNAME\username] FOR LOGIN [COMPUTERNAME\username]
END
GO

-- Grant database permissions
USE [master]
GO

IF NOT EXISTS (SELECT name FROM master.sys.database_principals WHERE name =
'db_role_ec2_sql_ha')
BEGIN
    CREATE ROLE [db_role_ec2_sql_ha]
END

GRANT VIEW DATABASE STATE to [db_role_ec2_sql_ha]
GO

ALTER ROLE [db_role_ec2_sql_ha] ADD MEMBER [COMPUTERNAME\username]
GO

-- Grant server permissions
USE [master]
GO

IF NOT EXISTS (SELECT name FROM master.sys.server_principals WHERE name =
'svr_role_ec2_sql_ha')
BEGIN
    CREATE SERVER ROLE [svr_role_ec2_sql_ha]
END

GRANT VIEW SERVER STATE TO [svr_role_ec2_sql_ha]
```

```
GRANT VIEW ANY DEFINITION TO [svr_role_ec2_sql_ha]
GRANT VIEW ANY DATABASE TO [svr_role_ec2_sql_ha]
GO

ALTER SERVER ROLE [svr_role_ec2_sql_ha] ADD MEMBER [COMPUTERNAME\username]
GO
```

Your local user setup is complete. When enabling Amazon EC2 instances for SQL HA standby detection, you can supply the ARN for the AWS secret you created.

## Disable Amazon EC2 High Availability for SQL Server

You can disable Amazon EC2 High Availability for SQL Server (SQL HA). Note only instances enabled by SQL HA can receive the SQL Server license savings. Use one of the following methods to disable SQL HA for your instances:

### Console

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. In the navigation panel, choose **Instances**.
3. Select the instances in the High Availability deployment to enable SQL HA standby detection monitoring, choose **Actions, Instance settings, Modify SQL High Availability settings**.
4. In the **Review prerequisites** step, choose **Next**. The prerequisites only apply for enabling the monitoring, and it is not necessary to review them for disabling SQL HA standby detection monitoring.
5. In the **Manage SQL High Availability license savings** step, for each instance to disable, for **SQL High Availability license savings**, select **None**.
6. Choose **Next**.
7. In the **Review and apply changes** step, review the configuration and then choose **Apply changes**.

### AWS CLI

Use the [disable-instance-sql-ha-standby-detections](#) command. For `instance-ids`, specify the IDs of the instances to disable.

```
aws ec2 disable-instance-sql-ha-standby-detections \  
--instance-ids instance_ids
```

## View states for Amazon EC2 High Availability for SQL Server

You can view the Amazon EC2 High Availability for SQL Server (SQL HA) current and historical states. Use one of the following methods:

### Console

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. In the navigation panel, choose **Instances**.
3. Select the instances in the High Availability deployment for which to view the SQL HA states, then choose the **SQL High Availability** tab.

### AWS CLI

To view the current SQL HA states for Amazon EC2 instances, use the [describe-instance-sql-ha-states](#) command. This command only shows the current SQL HA status of your onboarded instances.

```
aws ec2 describe-instance-sql-ha-states \  
--instance-ids instance_ids
```

To view the historical SQL HA states for instances, use the [describe-instance-sql-ha-history-states](#) command. This command returns your SQL HA instance state transitions in descending time order.

```
aws ec2 describe-instance-sql-ha-history-states \  
--instance-ids instance_ids \  
--start-time period_start_timestamp \  
--end-time period_end_timestamp
```

# Find a SQL Server license-included AMI

This topic describes how you can find SQL Server license-included AMIs using the Amazon EC2 console, the AWS Tools for PowerShell, the AWS CLI, or by searching the AWS Marketplace. For SQL Server license-included AMIs, use the installation and setup media included in C:\SQLServerSetup to make changes to the default installation, add new features, or install additional named instances.

As you select a SQL Server license-included AMI, consider the following requirements you might have for the instances that you'll launch:

- The AWS Region
- The operating system
- The architecture: 64-bit (x86\_64)
- The [root device](#) type: Amazon EBS-backed (EBS)
- The provider (for example, Amazon Web Services)
- Additional software (for example, SQL Server)

## Note

To view changes to each release of the AWS Windows AMIs, including SQL Server updates, see the [AWS Windows AMI version history](#) in the *Amazon EC2 User Guide*.

## Methods to find a SQL Server license-included AMI

### AWS Marketplace

To view a list of SQL Server AMIs available from AWS in AWS Marketplace, see [Windows AMIs](#).

### Console

You can find SQL Server license-included AMIs using the Amazon EC2 console. You can select from the list of AMIs when you use the launch instance wizard to launch an instance, or you can search through all available AMIs using the **Images** page. AMI IDs are unique to each AWS Region.

## To find a SQL Server license-included AMI using the launch instance wizard

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. From the navigation bar, select the Region in which to launch your instances. You can select any Region that's available to you, regardless of your location.
3. From the console dashboard, choose **Launch instances**.
4. Under **Application and OS Images (Amazon Machine Image)**, enter SQL in the search bar and choose **Enter**. You will be taken to the **AMIs** page, where you can browse and choose from AMIs with SQL Server included. You can choose from AMIs under the **Quickstart AMIs**, **My AMIs**, **AWS Marketplace AMIs**, and the **Community AMIs** tabs. You can filter by cost, operating system, and architecture.
5. To launch an instance from this AMI, select it and then choose **Launch instance**. For more information about launching an instance using the console, see [Launch an instance using the new launch instance wizard](#). If you're not ready to launch the instance now, take note of the AMI ID for later.

## To find a SQL Server AMI using the AMIs page

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. From the navigation bar, select the Region in which to launch your instances. You can select any Region that's available to you, regardless of your location.
3. In the navigation pane, choose **AMI Catalog**.
4. Enter SQL in the search bar and choose **Enter**. You can choose from SQL Server license-included AMIs under the **Quickstart AMIs**, **My AMIs**, **AWS Marketplace AMIs**, and the **Community AMIs** tabs. You can filter by cost, operating system, and architecture.
5. To launch an instance from this AMI, select it and then choose **Launch instance**. For more information about launching an instance using the console, see [Launching your instance from an AMI](#). If you're not ready to launch the instance now, take note of the AMI ID for later.

## PowerShell

You can use cmdlets for Amazon EC2 to list only the Windows AMIs that match your requirements. After locating an AMI that matches your requirements, take note of its ID so

that you can use it to launch instances. For more information, see [Launch an Instance Using Windows PowerShell](#) in the *AWS Tools for PowerShell User Guide*.

To list the latest SQL Server license-included AMIs provided by Amazon, you can use the `Get-SSMLatestEC2Image` cmdlet. The following command lists the latest Windows AMIs with *SQL* in their image name:

```
Get-SSMLatestEC2Image -Path ami-windows-latest -ImageName *SQL*
```

To list SQL Server license-included AMIs using commands that match specific criteria, you can use the `Get-EC2Image` cmdlet in addition to filters. The following commands filter for AMIs owned by you, or Amazon, with *SQL* in their name:

```
$name_values = New-Object 'collections.generic.list[string]'
$name_values.add("*SQL*")
$filter_name = New-Object Amazon.EC2.Model.Filter -Property @{Name = "name"; Values
    = $name_values}
Get-EC2Image -Owner amazon, self -Filter $filter_name
```

For more information and examples, see [Find an AMI Using Windows PowerShell](#) in the *AWS Tools for PowerShell User Guide*.

## AWS CLI

You can use AWS CLI commands for Amazon EC2 to list only the SQL Server license-included AMIs that match your requirements. After locating an AMI that matches your requirements, take note of its ID so that you can use it to launch instances. For more information, see [Launching an Instance Using the AWS CLI](#) in the *AWS Command Line Interface User Guide*.

The [describe-images](#) command supports filtering parameters. For example, use the `--owners` parameter with `amazon` to display public AMIs owned by Amazon or `self` to list AMIs you own. You can specify multiple values for the `--owners` parameter as in the following example:

```
aws ec2 describe-images --owners self amazon
```

You can add the following filter to the previous command to display only SQL Server license-included AMIs:

```
--filters "Name=name,Values=*SQL*"
```

You can use the following filter with the command to display only AMIs backed by Amazon EBS:

```
--filters "Name=root-device-type,Values=ebs"
```

You can combine multiple filters together. For example, this command will list all AMIs owned by you or Amazon with *SQL* in the AMI name and the `--root-device-type` parameter as `ebs`:

```
aws ec2 describe-images --owners self amazon --filters "Name=name,Values=*SQL*"  
"Name=root-device-type,Values=ebs"
```

 **Note**

Omitting the `--owners` flag from the `describe-images` command will return all images for which you have launch permissions, regardless of ownership.

# Deploy SQL Server on Amazon EC2

To launch Microsoft SQL Server using Amazon Elastic Compute Cloud (Amazon EC2) instances with Windows Server, perform the following steps according to your use case.

New SQL environment deployments are classified under three categories:

- SQL Server standalone
- SQL Server Failover Cluster Instances (FCI)
- SQL Server Always On availability groups (AG)

## Considerations

Before you launch SQL Server on your instance, consider the following:

- If you use an AWS provided AMI, you must initially manage SQL Server as the local administrator. For more information, see [Connect to SQL Server on Amazon EC2](#).
- If you use an instance type that provides an instance store, we recommend that you place tempdb on an instance store volume. This can improve performance and decrease costs. For more information, see [Place tempdb in an instance store](#).
- The built-in availability form of clustering in Windows Server is activated by a feature named Failover Clustering. This feature allows you to build a Windows Server Failover Cluster (WSFC) to use with an availability group or failover cluster instances (FCI).
- Always On is an umbrella term for the availability features in SQL Server, and the term covers both availability groups and FCIs. Always On isn't the name of the Always On availability group (AG) feature.
- The major difference between FCI and AG is that all FCIs require some sort of shared storage, even if it's provided through networking. The FCI's resources can be run and owned by one node at any given time. AG doesn't require that shared storage is also highly available. It's a best practice to have replicas that are local in one data center for high availability, and remote ones in other data centers for disaster recovery, each with separate storage.
- An availability group also has another component called the listener. The listener allows applications and end users to connect without needing to know which SQL Server instance is hosting the primary replica. Each availability group has its own listener.

# Deployment options

Use one of the following options to deploy SQL Server on Amazon EC2.

## Deploy SQL Server on Amazon EC2 with AWS Launch Wizard

AWS Launch Wizard is a service that guides you through the sizing, configuration, and deployment of enterprise applications following AWS Cloud best practices. AWS Launch Wizard for SQL Server supports both high availability and single instance deployments according to AWS and SQL Server best practices. For more information, see the [AWS Launch Wizard for SQL Server User Guide](#).

### Always On availability groups (AG)

Deploy your SQL Server Always On availability groups with primary and secondary replicas for database level protection. Each replica is hosted by a SQL Server instance with its own local storage.

### Always On Failover Cluster Instances (SQL FCI)

Deploy SQL Server Always On using Failover Cluster Instances (FCI) for instance-level protection. A single SQL Server instance is installed across Windows Server Failover Clustering (WSFC) nodes to ensure high availability and storage sharing.

Launch Wizard uses Amazon FSx to provide the following shared storage options required for SQL FCI deployments:

- Amazon FSx for NetApp ONTAP using Microsoft iSCSI endpoint
- Amazon FSx for Windows File Server using SMB 3.0 continuously available Windows file share

For more information on how to deploy SQL Server with Launch Wizard, see [Deploy an application with AWS Launch Wizard for SQL Server on Windows](#) in the *AWS Launch Wizard User Guide*.

## Deploy SQL Server standalone

For a SQL Server standalone deployment, you can use one of the license-include AMIs provided by AWS or by using your own licensed media. For a list of SQL Server AMIs provided by AWS, see [Windows AMIs](#). For more information on licensing options, see [Understand licensing options and considerations for Microsoft SQL Server on Amazon EC2](#).

## Deploy SQL Server failover cluster instances (FCIs)

Failover cluster instances (FCIs) provide availability for the entire installation of SQL Server known as an instance. Everything that is included in the instance, such as databases, SQL Server Agent jobs, and linked servers, move to a different server when the underlying server fails.

You can use AWS Launch Wizard to deploy SQL Server FCIs in the AWS Cloud. Launch Wizard identifies the AWS resources to automatically provision the SQL Server databases based on your use case. For more information, see [Get started with AWS Launch Wizard for SQL Server](#).

You can reference the following AWS blogs to manually deploy SQL Server FCIs:

- (Amazon FSx) [Deploy a SQL Server FCI using SMB 3.0 Continuously Available File Shares \(CAFS\) as shared storage](#)
- (Amazon FSx) [Deploy a SQL Server FCI using Microsoft iSCSI Initiator as shared storage](#)
- (Amazon EBS) [Deploy SQL Server FCI using Amazon EBS Multi-Attach with Persistent Reservations \(MAPR\)](#)

## Deploy SQL Server Always On availability groups (AG)

Always on availability groups provide high availability and disaster recovery of user databases through data replication. Availability groups can also distribute read operations amongst member nodes.

You can use [AWS Launch Wizard](#) to deploy a SQL Server Always On availability group in the AWS Cloud. Launch Wizard identifies the AWS resources to automatically provision the SQL Server databases based on your use case. For more information, see [Get started with AWS Launch Wizard for SQL Server](#).

To manually deploy a SQL Server Always On availability group, perform the following steps:

### Prerequisites

Before you manually deploy a SQL Server Always On availability group, you must perform the following prerequisites.

- Launch two Amazon EC2 instances with Windows Server 2016 or later and SQL Server 2016 or later Enterprise edition across two Availability Zones within an Amazon VPC. If the deployment is for testing purposes only, you can consider using SQL Server Developer edition.

- Configure secondary Amazon EBS volumes to host SQL Server Master Data File (MDF), Log Data File (LDF), and SQL Backup files (.bak). For more information on the volume types that you can use, see [Amazon EBS volume types](#) in the *Amazon Elastic Compute Cloud User Guide*.
- Deploy the cluster nodes in private subnets. You can then use Remote Desktop Protocol (RDP) to connect from a jump server to the cluster node instances.
- Configure inbound security group rules and [Windows firewall exceptions](#) to allow the nodes to communicate in a restrictive environment.
- Open all necessary ports for Active Directory domain controllers so that the SQL nodes and witness can join the domain and authenticate against Active Directory. For more information, see [Active Directory and Active Directory Domain Services Port Requirements](#) in the Microsoft documentation.
- Join the nodes to the domain before you create the Windows failover cluster. Ensure that you are logged in with domain credentials before you create and configure the cluster.
- Run the SQL Database instances with an Active Directory service account.
- Create a SQL login with sysadmin permissions using Windows domain authentication. Consult with your database administrator for details. For more information, see [Create a login using SSMS for SQL Server](#) in the Microsoft documentation.
- Properly configure the SQL browser for SQL Server named instances.

## Configure the secondary IPs for each cluster node elastic network interface

Two secondary IP addresses are required for each cluster node elastic network interface.

### Note

If you do not plan to deploy a SQL Group Listener, add only one secondary IP address for each cluster node elastic network interface.

1. Navigate to the [Amazon EC2 console](#) and choose the AWS Region where you want to host your Always On cluster.
2. Choose **Instances** from the left navigation pane, and then select your Amazon EC2 cluster instance.
3. Choose the **Networking** tab.

4. Under **Network interfaces**, select the network interface and then choose **Actions > Manage IP addresses**.
5. Choose the network interface Id to open the expandable section, and then choose **Assign new IP address**. You can enter a specific IP address or keep the default entry as Auto-assign. Repeat this step to add a second new IP address.
6. Choose **Save > Confirm**.
7. Repeat steps 1 through 7 for the other Amazon EC2 instance that will be included in the cluster.

## Create a two-node Windows cluster

Perform the following steps to create a two-node Windows cluster.

1. [Connect to your Amazon EC2 instance](#) with RDP, using a domain account with local administrator permissions on both nodes.
2. On the Windows **Start** menu, open **Control Panel**, and then choose **Network and Internet > Network and Sharing Center**.
3. Choose **Change adapter settings** from the left navigation pane.
4. Choose your network connection, and then choose **Change settings of this connection**.
5. Choose **Internet Protocol Version 4 TCP/IPv4**, and then choose **Properties**.
6. Choose **Advanced**.
7. Under the **DNS** tab, choose **Append primary and connection specific DNS suffixes**.
8. Choose **OK > OK > Close**.
9. Repeat steps 1 through 8 for the other Amazon EC2 instance to include in the cluster.
10. On each instance, install the cluster feature on the nodes from the Server Manager, or run the following Windows PowerShell command:

```
Install-WindowsFeature -Name Failover-Clustering -IncludeManagementTools
```

11. Open the command line as an administrator and enter `cluadmin.msc` to open the Cluster Manager.
12. Open the context menu (right-click) for **Failover Cluster Manager**, and then choose **Create Cluster**.
13. Choose **Next > Browse**.

14. For **Enter the object names to select**, enter the cluster node hostnames, and then choose **OK**.
15. Choose **Next**. You can now choose whether to validate the cluster. We recommend that you perform a cluster validation. If the cluster does not pass validation Microsoft may be unable to provide technical support for your SQL cluster. Choose **Yes** or **No**, and then choose **Next**.
16. Enter a **Cluster Name**, and then choose **Next**.
17. Clear **Add all eligible storage to the cluster**, and then choose **Next**.
18. When the cluster creation is complete, choose **Finish**.

 **Note**

Cluster logs and reports are located at %systemroot%\cluster\reports.

19. In the **Cluster Core Resources** section of Cluster Manager, expand the entry for your new cluster.
20. Open the context menu (right-click) for the first IP address entry, and then choose **Properties**. For **IP Address**, choose **Static IP Address**, and then enter one of the secondary IP addresses associated with the eth0 elastic network interface. Choose **OK**. Repeat this step for the second IP address entry.
21. Open the context menu (right-click) for the cluster name, and then choose **Bring Online**.

 **Note**

We recommend that you configure a [File Share Witness \(FSW\)](#) in addition to your cluster to act as a tie-breaker. You can also use [Amazon FSx for Windows File Server with Microsoft SQL Server](#).

## Create Always On availability groups

Perform the following steps to create Always On availability groups.

1. Open SQL Server Configuration Manager.
2. Open the context menu (right-click) for the SQL instance, and then choose **Properties**.
3. On the **AlwaysOn High Availability** tab, select **Enable AlwaysOn Availability Groups**, and then choose **Apply**.
4. Open the context menu (right-click) for the SQL instance, and then choose **Restart**.

5. Repeat steps 1 through 4 on the other cluster node to include in the cluster.
6. Open Microsoft SQL Server Management Studio (SSMS).
7. Log in to one of the SQL instances with your Windows authenticated login that has access to the SQL instance.

 **Note**

We recommend that you use the same MDF and LDF directory file paths across the SQL instances.

8. Create a test database. Open the context menu (right-click) for **Databases**, and then choose **New Database**.

 **Note**

Make sure that you use the **Full** [recovery model](#) on the **Options** page.

9. Enter a **Database name**, and then choose **OK**.
10. Open the context menu (right-click) for the new database name, choose **Tasks**, and then choose **Back Up For Backup type**, choose **Full**.
11. Choose **OK > OK**.
12. Open the context menu (right-click) for **Always On High Availability** and then choose **New Availability Group Wizard**.
13. Choose **Next**.
14. Enter an **Availability group name**, and then choose **Next**.
15. Select your database, and then choose **Next**.
16. A primary replica is already present in the Availability Replicas window. Choose **Add Replica** to create a secondary replica.
17. Enter a **Server name** for the secondary replica and then choose **Connect**.
18. [Decide which Availability Mode you want to use](#) for each replica, and then choose either **Synchronous commit** or **Asynchronous commit**.
19. Choose **Next**.
20. Choose your [data synchronization preference](#), and then choose **Next**.
21. When the validation is successful, choose **Next**.

**Note**

You can safely ignore **Checking the listener configuration** because you will add it later.

22. Choose **Finish** > **Close**.

## Add a SQL Group Listener

Perform the following steps to add a SQL Group Listener.

1. Open SQL Server Management Studio (SSMS) and expand **Always On High Availability, Availability Groups, <primary replica name>**.
2. Open the context menu (right-click) for **Availability Group Listeners** and then choose **Add Listener**. Enter a **DNS Name**.
3. Enter **Port** 1433.
4. Choose **Static IP** for **Network Mode**.
5. Choose **Add**.

For the **IPv4 Address**, enter the second secondary IP address from one of the cluster node instances, and then choose **OK**. Repeat this step using the second secondary IP address from the other cluster node instance.

6. Choose **OK**.

**Note**

If you receive errors when you add a SQL Group Listener, you may be missing permissions. For troubleshooting see:

- [Troubleshooting AlwaysOn availability group listener creation in SQL Server 2012](#)
- [Create Availability Group Listener Fails with Message 19471, 'The WSFC cluster could not bring the Network Name resource online'](#)

## Test failover

1. From SSMS, open the context menu (right-click) for the primary replica on the navigation menu, and then choose **Failover**.
2. Choose **Next > Next**.
3. Choose **Connect > Connect**.
4. Choose **Next**, and then choose **Finish**. The primary replica will become the secondary replica after failover.

# Connect to Microsoft SQL Server on Amazon EC2

You can connect to your Microsoft SQL Server instance using one of the following tools.

## Topics

- [SQL Server Management Studio \(SSMS\)](#)
- [SQL Server Configuration Manager](#)

## SQL Server Management Studio (SSMS)

By default, only the built-in local administrator account can access a SQL Server instance launched from an AWS Windows AMI. You can use SQL Server Management Studio (SSMS) to add domain users so that they can access and manage SQL Server.

Perform the following steps to access a SQL Server instance on Amazon EC2 as a domain user.

1. [Connect to your instance](#) as a local administrator using Remote Desktop Protocol (RDP).
2. Open SQL Server Management Studio (SSMS).
3. For **Authentication**, choose **Windows Authentication** to log in with the built-in local administrator.
4. (Optional) Allow domain users to log in.
  - a. Choose **Connect**.
  - b. In Object Explorer, expand **Security**.
  - c. Open the context menu (right-click) for **Logins** then select **New Login**.
  - d. For **Login name**, select **Windows authentication**. Enter **Domain\username**, replacing **DomainName** with your domain NetBIOS name and **username** with your Active Directory user name.
  - e. On the **Server roles** page, select the [server roles](#) that you want to grant to the Active Directory user.
  - f. Select the **General** page, and then choose **OK**.
  - g. Log out from the instance and then log in again as a domain user.
  - h. Open SSMS. For **Authentication**, choose **Windows authentication** to log in with your domain user account.

- i. Choose **Connect**.

## SQL Server Configuration Manager

To connect to SQL Server using SQL Server Configuration Manager, see [SQL Server Configuration Manager](#) in the Microsoft documentation.

# Microsoft SQL Server on Amazon EC2 backup with the AWS VSS solution

You can use the AWS VSS solution to create snapshots of the volumes attached to your EC2 instance, including data from your Microsoft SQL Server database. The AWS VSS solution uses Amazon managed command documents with the [AWS Systems Manager Run Command](#). The snapshot process uses the Windows [Volume Shadow Copy Service \(VSS\)](#) to create EBS volume-level backups without requiring you to shut down or pause your SQL Server application or disconnect active connections.

There is no additional cost to use the AWS VSS solution to create EBS snapshots. You only pay for EBS snapshots created by the backup process. For more information, see [How am I billed for my Amazon EBS snapshots?](#)

To restore your SQL Server databases from AWS VSS solution based EBS snapshots, see [Use an automation runbook to restore your database from AWS VSS solution snapshots](#).

## AWS VSS solution based snapshot prerequisites

To start backing up your SQL Server databases with AWS VSS solution based EBS snapshots, you must meet all of the prerequisites.

- Your instance must satisfy the system requirements, and have the appropriate IAM permissions attached to your instance profile role. If you use Systems Manager to run the recommended command document (AWSEC2-VssInstallAndSnapshot), you can skip the component install prerequisite. That command document automatically installs the VSS components as needed on your instance.

For more information, see [Prerequisites to create Windows VSS based EBS snapshots](#) in the *Amazon EC2 User Guide*.

- To enable the restore process for your SQL Server database, you must set your database to full recovery mode before you take the snapshot. To configure recovery mode, see [View or change the recovery model of a database \(SQL Server\)](#) on the *Microsoft Learn* website.

## Create AWS VSS solution based EBS snapshots

To create EBS snapshots that you can use to restore your database instance, see [Use Systems Manager command documents to create VSS based snapshots](#). We recommend that you run the **AWSEC2-VssInstallAndSnapshot** command document, which automatically installs the VSS components as needed on your instance.

### Note

To ensure that you can use the AWS VSS solution snapshots to restore your database, set the `SaveVssMetadata` parameter to `True` for your command document before you run it.

# Use an automation runbook to restore your database from AWS VSS solution snapshots

This guide explains how to use the `AWSEC2-RestoreSqlServerDatabaseWithVss` automation runbook to restore a Microsoft SQL Server database running on an EC2 instance from application-consistent snapshots created by the AWS VSS solution. You can tailor the restoration parameters to your specific needs, such as setting the database to restore mode after restoration. By leveraging Windows VSS technology, this solution offers the following advantages:

- Fast restoration times.
- The ability to perform the restore without shutting down or pausing the Microsoft SQL Server application.

For more information about the AWS VSS solution, see [Application consistent Windows VSS based Amazon EBS snapshots](#) in the *Amazon EC2 User Guide*.

## Contents

- [Pricing](#)
- [VSS based database restore solution change history](#)
- [VSS snapshot restore prerequisites](#)
- [Restore your SQL Server database from VSS snapshots](#)

## Pricing

The AWS VSS solution uses AWS Systems Manager automation runbooks with EBS resources to restore a Microsoft SQL Server database on an EC2 instance. Associated costs include the following:

### Systems Manager Automation

With Systems Manager automation runbooks, you pay only for what you use and are charged based on the number and duration of steps, which includes a free tier per account. If you created an organization, your free tier usage is shared across all accounts in the Consolidated Billing family. For more information about Systems Manager automation pricing, see [Automation](#).

## Amazon EBS

During restore steps, the AWS VSS solution creates a new EBS volume and restores data from VSS based EBS snapshots. With Amazon EBS resources, you pay only for what you provision. For more information, see [Amazon EBS pricing](#).

## VSS based database restore solution change history

The following table includes release and change history for the AWS VSS restore solution.

| Release date     | Details                                                                                                                                                                                                                  |
|------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| January 14, 2026 | Enhanced restore solution to create new Amazon EBS volumes only from snapshots that contain source database files to be restored, reducing cleanup overhead and improving the overall reliability of restore operations. |
| April 22, 2025   | Fixed an issue where the restore process times out when disk partition information is not available, improving reliability of restore operations.                                                                        |
| January 24, 2025 | <ul style="list-style-type: none"><li>• Updated automation document description to improve clarity</li><li>• Minor bug fixes and stability improvements</li></ul>                                                        |
| January 15, 2025 | Initial release of the VSS restore solution.                                                                                                                                                                             |

## VSS snapshot restore prerequisites

To restore your SQL Server databases from AWS VSS solution based EBS snapshots, you must meet the following prerequisites.

**Note**

The `AWSEC2-RestoreSqlServerDatabaseWithVss` automation runbook only supports restoring snapshots to the original EC2 instance where the snapshots were created.

- **Disk management configuration** – Your EC2 database instance must be configured with Basic Disks. For more information, see [Basic Disks](#) on the *Microsoft Learn* website.
- **Microsoft SQL Server deployment options** – To restore a SQL Server database with the `AWSEC2-RestoreSqlServerDatabaseWithVss` automation runbook, the database must either be configured as a standalone deployment, or be the primary database in a Microsoft SQL Server Always On availability group. For more information, see [Deployment options](#).
- **Configure settings to save VSS metadata files** – To successfully initiate a restore operation, VSS metadata files are required. The following files are generated for each snapshot set taken during the snapshotting process.
  - `{Snapshot set id}-{timestamp}-BCD.xml`
  - `{Snapshot set id}-{timestamp}-SqlServerWriter.xml`
  - `{Snapshot set id}-{timestamp}-VolumeMapping.json`

**Note**

The volume mapping metadata file (`{Snapshot set id}-{timestamp}-VolumeMapping.json`) maps Windows drives to their corresponding snapshots and is used in VSS restore operations to create EBS volumes from snapshots that contains database files to be restored.

To ensure that these files are generated, set the `SaveVssMetadata` parameter to `true` when you run the command document.

- [Grant IAM permissions for the restore process.](#)

## Grant IAM permissions for the restore process

Executing the `AWSEC2-RestoreSqlServerDatabaseWithVss` automation runbook to restore databases needs permissions to perform necessary Amazon EC2 and Systems Manager operations. Follow these steps to grant the appropriate permissions.

1. [Attach the `AWSEC2VssRestorePolicy` managed policy to the role that's used for the automation execution.](#)
2. [Grant IAM permissions to the invoker role for starting and managing automation executions.](#)

### Attach the `AWSEC2VssRestorePolicy` managed policy to the role that's used for the automation execution

You can choose from the following options to attach the **`AWSEC2VssRestorePolicy`** AWS managed policy to the role that Systems Manager uses for interacting with Amazon EC2 and Systems Manager when executing the `AWSEC2-RestoreSqlServerDatabaseWithVss` automation runbook. For more information about this managed policy, see [AWSEC2VssRestorePolicy](#).

- Create a role, attach the **`AWSEC2VssRestorePolicy`** managed policy, and add a `PassRole` policy to restrict access. Use the ARN of this role for the `AutomationAssumeRole` parameter when invoking the automation, and the automation execution will assume this role. Expand the `Invoke` automation with an assumed role (`recommended`) section to see detailed steps.
- Attach the **`AWSEC2VssRestorePolicy`** managed policy to the invoker role that initiates the automation execution, without specifying the `AutomationAssumeRole` parameter. For example, if you start the automation execution from the AWS console, the console role acts as the invoker role. Expand the `Invoke` automation without an assumed role section to see detailed steps.

### Invoke automation with an assumed role (recommended)

#### Step 1: Create the role that the automation assumes and attach your policy

1. Open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the navigation pane, choose `Roles`, and then choose **Create role**. This opens the **Select trusted entity** page.
3. In the **Trusted entity type** panel, choose **AWS service**. This is the default selection.

4. In the **Use case** panel, select **Systems Manager** from the list, and then choose **Next**. This opens the **Add permissions** page.
5. Search for **AWSEC2VssRestorePolicy**. Select the check box next to the name and then choose **Next**. This takes you to the **Name, review, and create** page.
6. In the **Role details** panel, enter **Role name** and **Description**.
7. When you've finished reviewing, choose **Create role**. This takes you back to the **Roles** page.
8. Open the detail page for the role that you just created. Take note of the **Role Name** at the top for future reference.

Copy the **Role ARN** from the **Summary** panel to use in the next steps, then continue to Step 2 to create a PassRole policy for your role.

## Step 2: Create an inline policy to pass the role that the automation assumes

1. In the detail page for the role that you just created, choose the **Permissions** tab.
2. Choose **Add inline policy** from the **Add permissions** menu. This opens the **Specify permissions** page.
3. Select the **Visual** policy editor.
4. Choose **IAM** from the **Service** list.
5. In the **Actions allowed** search box, enter `PassRole`, then select the **PassRole** check box.
6. The **Resources** panel opens with the **Specific** option selected by default. Select the **Add ARNs** link to open a panel where you can specify the ARN for your role.
7. In the **Resource ARN** box, paste the ARN that you copied at the end of Step 1. IAM automatically populates the role name based on the ARN.
8. Choose **Add ARNs** to save your resource ARN. This takes you back to the **Specify permissions** page, and shows your entry.
9. Choose **Next** to review your policy. This opens the **Review and create** page.
10. On the Review Policy page, enter a name (for example, `VssRestorePassRolePolicy`) and then choose **Next** to create the PassRole policy for your role.

## Invoke automation without an assumed role

1. Open the IAM console at <https://console.aws.amazon.com/iam/>.

2. In the navigation pane, choose **Roles**, and then select the role that will be used to start the automation execution. For example, if you will start the automation execution from console, you should choose the current console role, which appears in the upper right corner of the console:

```
role/user @ account
```

3. In the **Permissions** tab, choose **Attach policies** from the **Add permissions** menu. This opens the **Attach policy to <selected role>** page.
4. Use the search bar in the **Other permissions policies** panel to search for **AWSEC2VssRestorePolicy**. Select the check box next to the name and then choose **Add permissions**.

## Grant IAM permissions to the invoker role for starting and managing automation executions

To attach necessary permissions to the role that starts and manages the AWSEC2-RestoreSqlServerDatabaseWithVss automation executions, follow these steps.

1. Open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the navigation pane, choose **Roles**, and then select the role that will be used to start the automation execution.
3. Choose **Add inline policy** from the **Add permissions** menu. This opens the **Specify permissions** page.
4. Select the **JSON** policy editor and copy the following JSON policy content into the editor. The policy allows the role to:
  - Execute the AWSEC2-RestoreSqlServerDatabaseWithVss automation runbook.
  - Stop and send signals to an automation execution.
  - View details about the automation execution after it has been started.

### JSON

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Effect": "Allow",  
      "Action": "aws:runbook:execute",  
      "Resource": "*"
```

```

    {
      "Sid": "StartVssRestoreAutomationExecution",
      "Effect": "Allow",
      "Action": "ssm:StartAutomationExecution",
      "Resource": [
        "arn:aws:ssm:*:*:document/AWSEC2-
RestoreSqlServerDatabaseWithVss",
        "arn:aws:ssm:*:*:automation-execution/*"
      ]
    },
    {
      "Sid": "ManageVssRestoreAutomationExecution",
      "Effect": "Allow",
      "Action": [
        "ssm:StopAutomationExecution",
        "ssm:GetAutomationExecution",
        "ssm:DescribeAutomationExecutions",
        "ssm:DescribeAutomationStepExecutions",
        "ssm:SendAutomationSignal"
      ],
      "Resource": [
        "arn:aws:ssm:*:*:automation-execution/*"
      ]
    }
  ]
}

```

5. If you are to start the AWSEC2-RestoreSqlServerDatabaseWithVss automation with an assume role by providing a role arn to the AutomationAssumeRole parameter, you will need to add the following permission to the above policy statements, and replace the [AutomationAssumeRole's ARN] placeholder with the ARN of the role created in step Invoke runbook automation with an assumed role (recommended). The permission allows the invoker role to pass the automation assume role to Systems Manager.

```

{
  "Action": "iam:PassRole",
  "Effect": "Allow",
  "Resource": [
    "[AutomationAssumeRole's ARN]"
  ]
}

```

6. Choose **Next** to review your policy. This opens the review and create page.

7. On the **Review Policy** page, enter a name (for example, `VssRestoreRunSSMAutomationPolicy`) and then choose **Next** to create and add the inline policy to your role.

## Restore your SQL Server database from VSS snapshots

The `AWSEC2-RestoreSqlServerDatabaseWithVss` automation runbook provides a streamlined process to restore your SQL Server databases. This guide outlines the automation runbook functionality and explains the parameters that you can customize to suit your specific restoration needs.

Before you run the `AWSEC2-RestoreSqlServerDatabaseWithVss` automation runbook, ensure that you've met all prerequisites to create application consistent snapshots with the AWS VSS solution. For more information, see [Prerequisites to create Windows VSS based EBS snapshots](#) in the *Amazon EC2 User Guide*.

The `AWSEC2-RestoreSqlServerDatabaseWithVss` process consists of several key steps, as follows.

1. The first step uses `AWS-ConfigureAwsPackage` to upgrade or install the latest version of the `AwsVssComponents` component package.
2. The next step invokes `AWSEC2-PrepareVssRestore` to verify that prerequisites are met and that the input parameters include a valid value for the VSS Snapshot Set ID and Source Database Name.
3. The process then creates new EBS volumes from the snapshots and attaches them to the instance.
4. Finally, the process invokes `AWSEC2-RunVssRestoreForSqlDatabase`, which runs the Amazon EC2 VSS Agent to restore the database on the instance, and returns volume IDs and their usage status by the restored database, the final restore operation status, and Amazon EC2 VSS Agent logs.

## Parameters for the SQL Server database restore runbook

The `AWSEC2-RestoreSqlServerDatabaseWithVss` automation runbook uses the following input parameters:

**Note**

You can provide one of the following parameters to use a specific snapshot:

- `SnapshotSetId`
- `RestorePointOfTime`

If both parameters are empty, the restore uses the most recent snapshot set.

**InstanceId** (string, required)

The ID of the Amazon EC2 instance where the restore is performed.

**SourceDatabaseName** (string, required)

The name of the database that's included in the snapshots.

**TargetDatabaseName** (string, optional)

The restore process creates a new database, and restores the data from the snapshots to the new database from the snapshots. You can optionally set the name, or leave this parameter empty to use the default name for the new database (`Db_Restored`). The old database files are removed from the volume after the process completes.

**SnapshotSetId** (string, optional)

The Snapshot Set ID of the snapshot to use for recovery.

**RestorePointOfTime** (string, optional)

If this parameter is specified, the restore process uses the last Snapshot Set that was created before the provided point in time value. This parameter uses the following string format: **MM-dd-yyyy:hh-mm**.

**RestoreWithNorecovery** (string, required)

If this parameter is set to `True` the restore process leaves the database in restoring state so that you can apply transaction logs after the database restore is completed. To bring the database online immediately after the restore is completed, set this parameter to `False`.

**MetadataPath** (string, optional)

The fully qualified path to the directory where the VSS metadata files are stored. If not specified, the system uses the following default location, where metadata files are automatically saved during snapshot operations. Use this parameter to indicate a custom storage location if you've relocated the files. %PROGRAMDATA%\Amazon\AwsVss\VssMetadata.

**AutomationAssumeRole** (string, conditional)

The ARN of the IAM role that the automation assumes during execution. If not specified, the automation uses the IAM role that initiated the execution. For example, when starting the automation from the AWS Console without specifying this parameter, the automation uses your current console session's IAM role to interact with Amazon EC2 and SSM.

**ExecutionTimeout** (string, optional)

The amount of time, in seconds, that the RunVssRestoreForSqlDatabase step can run before it fails. If this value is not specified, the default timeout is 600 seconds.

## Run the SQL Server database restore process

### 1. Always On databases: Remove the source database from the SQL Server availability group

If your database is the primary database in an Always On availability group, you must remove the database from the availability group before you run the restore process.

- To remove the database from the availability group, follow the steps described in [Remove a primary database from an Always On availability group](#) on the *Microsoft Learn* website.
- Verify that the database remains online, and is not in a Synchronized state.

### 2. Execute AWSEC2-RestoreSqlServerDatabaseWithVss Automation Runbook

To view instructions, select the tab that matches your environment.

#### AWS Management Console

To run the restore in the AWS Management Console, follow these steps:

- Open the AWS Systems Manager console at <https://console.aws.amazon.com/systems-manager/>.

2. Select **Automation** from the navigation pane, under **Change Management Tools**. This shows a list of automation executions in your account, if applicable.
3. Choose **Execute automation**. This opens the **Choose runbook** page.
4. In the **Owned by Amazon** tab, search for `AWSEC2-RestoreSqlServerDatabaseWithVss`, and select it from the results. This opens the **Runbook details** panel.
5. Select **Default version at runtime** from the **Runbook version** list.
6. Choose **Next**. This opens the page where you can configure the settings and enter input parameters for the runbook.
7. Enter values for the **Input parameters** to configure runtime settings for the restore process. For parameter details, see [Parameters for the SQL Server database restore runbook](#).
8. Choose **Execute** to run the automation.

To review the execution status, navigate to the **Executed Steps** section within the automation execution details. This section displays all of the steps that ran, along with their runtime status. If the automation execution failed, follow the troubleshooting steps outlined in [Troubleshoot restoring your SQL Server database from AWS VSS solution snapshots using the Systems Manager console](#).

- Locate the command execution ID in the step details.
- Select the linked ID to access the execution details.
- Inspect the command output and return code for further troubleshooting.

## AWS CLI

Run the following command to restore a Microsoft SQL Server database on an instance. Replace or add parameters based on your specific use case. For parameter details, see [Parameters for the SQL Server database restore runbook](#).

```
aws ssm start-automation-execution \  
  --document-name "AWSEC2-RestoreSqlServerDatabaseWithVss" \  
  --parameters \  
  '{"InstanceId":"i-1234567890abcdef0","SourceDatabaseName":"DB_Source","TargetDatabaseName":'}
```

## Get execution status

To get the status of the automation execution, run the following command using the execution ID returned from `start-automation-execution`.

```
aws ssm get-automation-execution \
  --automation-execution-id a1b2c3d4-5678-90ab-cdef-EXAMPLE11111
```

## PowerShell

Run the following PowerShell commands with AWS Tools for Windows PowerShell to restore a SQL Server database on an instance. Replace or add parameters based on your specific use case. For parameter details, see [Parameters for the SQL Server database restore runbook](#).

```
Start-SSMAutomationExecution `
-DocumentName "AWSEC2-RestoreSqlServerDatabaseWithVss" `
-Parameter @{"InstanceId" = @($InstanceId); "SourceDatabaseName"
= @("DB_Source"); "TargetDatabaseName" = @("DB_Restored");
"RestoreWithNorecovery" = @($True); "AutomationAssumeRole" = @("Arn:of:role")}
```

## Get execution status

To get the status of the automation execution and the status of each action step, run the following command using the execution ID returned from `Start-SSMAutomationExecution`.

```
Get-SSMAutomationExecution -AutomationExecutionId $ExecutionId | Select-Object -
ExpandProperty StepExecutions | Select-Object -Property StepName, StepStatus |
Out-String
```

### 3. (Optional) Clean up unused EBS volumes after the automation execution succeeds

When volume mapping metadata is used, the `AWSEC2-RestoreSqlServerDatabaseWithVss` automation creates new EBS volumes only from volume snapshots that contain files of the database to be restored. However, when this metadata is not used, the automation creates a new EBS volume for each volume snapshot within the VSS snapshot set and attaches them to the target instance.

To determine if volume mapping metadata was used for your database restore operation, check the automation execution steps. In the automation execution details, examine the `PrepareForVolumeCreation` step output. If `ExecutionLog` in the `OutputPayload` says `No volume mapping found - using all snapshots`, volume mapping metadata was not used during the restore operation.

If volume mapping metadata was not used for your restore operation, follow these steps to identify and clean up volumes that don't contain restored database files:

- a. In the **Execution detail** page, choose the **RunVssRestoreForSqlDatabase** step (this is the last step).
- b. Choose the **CommandId** link in the **Outputs** section, and then choose the instance id to view the run command output.
- c. At the end of the output is a list of all volumes created and attached to the instance for restore purposes, and the status for each one. The status is either `in-use` or `unused`. To detach and delete the volumes, see [Detach an Amazon EBS volume from an Amazon EC2 instance](#) in the *Amazon EBS User Guide*.

## Troubleshoot restoring your SQL Server database from AWS VSS solution snapshots using the Systems Manager console

Before you try any troubleshooting steps, we recommend that you verify that you've met all [VSS snapshot restore prerequisites](#).

In the Systems Manager console, the **Failure details** section in the automation runbook **Execution Details** page includes the following information:

- **Failure Message**
- **Failure Type**
- **Failure Stage**

Together, these details offer a general overview of the cause of the failure. For a more comprehensive understanding, you must examine the specific step execution that failed.

The steps in `AWSEC2-RestoreSqlServerDatabaseWithVss` can be classified into three main categories:

## Script Execution Steps

These use the `aws:executeScript` action and include the following steps.

- `ExtractPrepareVssRestoreOutput`
- `PrepareForVolumeCreation`
- `ExtractCurrIterValues`
- `ConcatVolumeIds`

If one of these steps fails, investigate the step execution and review the execution logs found under **OutputPayload** in the **Outputs** section to determine what caused the issue.

## EC2 API Interaction Steps

These interact with Amazon EC2 APIs to create volumes from snapshots, attach them to instances, and monitor volume status. Many of these steps are performed within loop steps. If a loop step fails, identify the specific step within the loop that caused the failure to pinpoint the root cause. The **Failure details** section in the step execution details page provides relevant information for debugging.

## Run Command Execution Steps

These use the `aws:RunCommand` action to execute commands on the target instance. If a failure occurs due to a run command execution, examine the step execution details. Under **Outputs**, select the **CommandId** link for the command to access the **Run command execution** page, where you can view the complete log for debugging purposes.

# Evaluate if you can downgrade your Microsoft SQL Server edition

If you find that you aren't using Enterprise edition features, you can consider downgrading to Microsoft SQL Server Standard or Developer edition. By downgrading the edition, you can save on licensing costs.

## Note

SQL Server Developer edition is only eligible for use in non-production, development, and test workloads.

## Downgrade requirements

Your Microsoft SQL Server on Amazon EC2 must use the Bring Your Own License model (BYOL) and SQL Server Enterprise edition to be eligible for an in-place downgrade. If your instance meets this criteria, you should carefully evaluate which features are being used on your SQL Server instance before performing any changes. You can review the following SQL Server Enterprise edition features and instance level constraints to help evaluate your downgrade eligibility.

## Tip

A script is available to help evaluate if you can downgrade your SQL Server edition. For more information, see [Downgrade SQL Server Enterprise edition using AWS Systems Manager Document to reduce cost](#).

Confirm that your instance has **less** than the following resources available:


- 48 vCPUs
- 128 GiB of memory

If your instance is under-utilized for your workload, you can change the instance type or size your instance to meet these requirements. For more information, see [Change the instance type](#).

Confirm that you aren't using any of the following SQL Server Enterprise edition features:

- Database-level enterprise features
- Always On availability groups
- Online index operations
- Resource Governor
- Peer-to-peer or Oracle replication
- R or Python extensions
- Memory-optimized tempdb metadata

The following diagram walks through an evaluation for some of the downgrade requirements:



If your workload doesn't utilize any of the previously listed features, you should continue to evaluate if you use any less common SQL Server Enterprise edition features. For more information about SQL Server Enterprise editions and supported features, see the Microsoft documentation for your SQL Server version:

- [SQL Server 2025](#)
- [SQL Server 2022](#)
- [SQL Server 2019](#)
- [SQL Server 2017](#)
- [SQL Server 2016](#)
- [SQL Server 2014](#)

# Downgrade your SQL Server Enterprise edition

If you determine that you can downgrade your SQL Server Enterprise edition, you can follow this process to convert to SQL Server Standard or Developer edition. For information on how to automate for this process, see [Downgrade SQL Server Enterprise edition using AWS Systems Manager Document to reduce cost](#).

## Important

- This process will require downtime for your SQL Server instance. Your database will not be operational until the entire procedure has been completed successfully.
- Only SQL Server instances using BYOL software support in-place downgrading. For more information, see [Licensing options](#).

## To downgrade your SQL Server Enterprise edition

1. [Create a Full backup](#) of all user and system databases. Ensure that the backup completes successfully before continuing.
2. Note your current SQL Server minor version, service pack, cumulative updates, and the General Distribution Release (GDR). For more information, see [Determine which version and edition of SQL Server Database Engine is running](#) in the Microsoft documentation.
3. [Detach](#) all user databases.
4. [Stop the SQL Server Database Engine](#) service and copy the log and system database data files —master, model, and msdb—to a local backup folder.
5. [Uninstall SQL Server](#) Enterprise edition including all components.
6. [Reboot](#) the instance.
7. [Install SQL Server](#) Standard or Developer edition according to your requirement.
8. Install the same service packs and cumulative updates that you had before the uninstall.
9. Stop the SQL Server Database Engine service.
10. Using the backups you made in step 4, restore the master, model, and msdb databases.
11. Start SQL Server service.
12. [Attach](#) the mdf and ldf user databases that were detached in step 3 to your SQL Server instance.

13. Confirm that your database is operating as expected.

# Migrating an on-premises database to Amazon EC2

You can migrate your on-premises Microsoft SQL Server database to Amazon Elastic Compute Cloud (Amazon EC2). If you select a migration method and perform these steps, your on-premises database will reside on an Amazon EC2 instance running Windows Server.

## On-premises migration methods

- [Automated SQL Server backup and restore](#)
- [Manual SQL Server backup and restore](#)
- [Server rehost](#)

## Automated SQL Server backup and restore

You can use AWS Migration Hub Orchestrator to orchestrate and automate the migration of SQL Server databases to Amazon EC2 using automated native backup and restore. This feature of AWS Migration Hub uses predefined workflow templates that are built based on best practices. Migration Hub Orchestrator automates error-prone manual tasks involved in the migration process, such as checking environment readiness and connections. For more information, see [Rehost SQL Server on Amazon EC2](#) in the *Migration Hub Orchestrator User Guide*.

## Manual SQL Server backup and restore

You can use native backup files as a way to restore SQL Server databases without additional dependencies. You can back up and restore individual databases, or the entire database instance, from on premises to your EC2 instance.

### Manual migration topics

- [Prerequisites](#)
- [Step 1: Backing up your database](#)
- [Step 2: Uploading your database backup files](#)
- [Step 3: Downloading your database backup files](#)
- [Step 4: Restoring your database backup files](#)

## Prerequisites

You must meet the following prerequisites to migrate an on-premises database to Amazon EC2 using Amazon Simple Storage Service (Amazon S3):

- An active AWS account. For more information, see [Set up Microsoft SQL Server on Amazon EC2](#).
- A source SQL Server database running on premises that you'd like to migrate.
- A destination EC2 instance running Windows Server with SQL Server installed on it. It is preferred that the destination instance's SQL Server version is the same or higher than the source SQL Server version running on premises. For more information on how to launch an instance, see [Launch your instance](#) in the *Amazon EC2 User Guide*.
- An Amazon Simple Storage Service (Amazon S3) bucket. For more information, see [Creating, configuring, and working with Amazon S3 buckets](#) in the *Amazon S3 User Guide*.
- Microsoft SQL Server Management Studio (SSMS) has been installed on the destination EC2 instance. For more information, see [Download SQL Server Management Studio \(SSMS\)](#) in the Microsoft documentation.

## Step 1: Backing up your database

You will need to create a full backup of the database as well as back up the Transaction Log for the on-premises SQL Server to capture all of the necessary data for restoration. This procedure generates the backup files can restore your database with in an EC2 instance.

### To back up an on-premises database

1. Create a full backup of your database. For more information about how to create a full backup of your database, see [Create a Full Database Backup](#) in the Microsoft documentation.
2. Create a backup of the Transaction Log. For more information about how to back up the transaction log, see [Back up a Transaction Log](#) in the Microsoft documentation.
3. Make a note of the backup file locations, because you will need to upload them to Amazon S3 in the next step.

## Step 2: Uploading your database backup files

With the backup files created, you can now upload them to Amazon S3.

## To upload your database backup files

1. Determine size of your backup files to see which upload methods are supported.
2. Use the file locations you noted previously to upload your backup files. For more information about how you can upload your database backup files to Amazon S3, see [Uploading objects](#).

## Step 3: Downloading your database backup files

Once the backup files have been uploaded to Amazon S3, you can restore them in an EC2 instance.

### To download your backup files from Amazon S3 in the EC2 instance

1. Connect to your SQL Server instance and open SSMS. For more information, see [Connect to Microsoft SQL Server on Amazon EC2](#).
2. Download the backup files in your Amazon EC2 instance running SQL Server. For more information about downloading your files from Amazon S3, see [Downloading an object](#).
3. Make a note of the backup file locations, because you will need them to restore the database in the next step.

## Step 4: Restoring your database backup files

After you download the backup files, you can connect to your instance and restore them using SSMS.

### To restore your database

1. Connect to your instance and open SSMS.
2. Restore the full database backup using the backup files noted previously. For more information about restoring your database from the backup files, see [Restore a Database Backup Using SSMS](#) in the Microsoft documentation.
3. In the EC2 instance, validate that your database has been restored as expected.

## Server rehost

You can choose to *rehost* (lift and shift) your entire SQL Server to Amazon EC2 instead of individual databases using AWS Transform MGN or AWS Migration Hub Orchestrator.

## MGN (MGN)

MGN (MGN) automates the migration of your servers and applications to the cloud during a cutover window. For more information on how you can rehost SQL Server using MGN, see [Quick start guide](#) in the *MGN User Guide*.

## Migration Hub Orchestrator

Migration Hub Orchestrator orchestrates and further automates the rehost process for servers and applications. For more information on how you can rehost SQL Server using Migration Hub Orchestrator, see [Rehost applications on Amazon EC2](#) in the *Migration Hub Orchestrator User Guide*.

# Windows to Linux replatforming assistant for Microsoft SQL Server Databases

The Windows to Linux replatforming assistant for Microsoft SQL Server Databases service is a scripting tool. It helps you move existing Microsoft SQL Server workloads from a Windows to a Linux operating system. You can use the replatforming assistant with any Windows Server virtual machines (VMs) hosted in the cloud, or with on-premises environments running Microsoft SQL Server 2008 and later. The tool checks for common incompatibilities, exports databases from the Windows VM, and imports into an EC2 instance running Microsoft SQL Server 2017 on Ubuntu 16.04. The automated process results in a ready-to-use Linux VM configured with your selected SQL Server databases that can be used for experimenting and testing.

## Contents

- [Concepts](#)
- [Related services](#)
- [How Windows to Linux replatforming assistant for Microsoft SQL Server works](#)
- [Components](#)
- [Replatforming script prerequisites](#)
- [Run the Windows to Linux replatforming assistant for SQL Server script](#)

## Concepts

The following terminology and concepts are central to your understanding and use of the Windows to Linux replatforming assistant for Microsoft SQL Server Databases.

### Backup

A Microsoft SQL Server backup copies data or log records from a Microsoft SQL Server database or its transaction log to a backup device, such as a disk. For more information, see [Backup Overview \(Microsoft SQL Server\)](#).

### Restore

A logical and meaningful sequence for restoring a set of Microsoft SQL Server backups. For more information, see [Restore and recovery overview \( SQL Server\)](#).

## Replatform

A Microsoft SQL Server database can be replatformed from an EC2 Windows instance to an EC2 Linux instance running Microsoft SQL Server. It can also be replatformed to the VMware Cloud running Microsoft SQL Server Linux on AWS.

## Related services

[AWS Systems Manager \(Systems Manager\)](#) gives you visibility and control of your infrastructure on AWS. The Windows to Linux replatforming assistant for Microsoft SQL Server Databases uses Systems Manager to move your Microsoft SQL databases to Microsoft SQL Server on EC2 Linux. For more information about Systems Manager, see the [AWS Systems Manager User Guide](#).

## How Windows to Linux replatforming assistant for Microsoft SQL Server works

Windows to Linux replatforming assistant for Microsoft SQL Server Databases allows you to migrate your Microsoft SQL Server databases from an on-premises environment or from an EC2 Windows instance to Microsoft SQL Server 2017 on EC2 Linux using backup and restore. For the destination EC2 Linux instance, you provide either the EC2 instance ID or the EC2 instance type with the subnet ID and EC2 Key Pair.

When you run the PowerShell script for the Windows to Linux replatforming assistant for Microsoft SQL Server Databases on the source Microsoft SQL Server databases, the Windows instance backs up the databases to an encrypted [Amazon Simple Storage Service \(S3\)](#) storage bucket. It then restores the backups to an existing Microsoft SQL Server on EC2 Linux instance, or it launches a new Microsoft SQL Server on EC2 Linux instance and restores the backups to the newly created instance. This process can be used to replatform your 2-tier databases running enterprise applications. It also enables you to replicate your database to Microsoft SQL Server on Linux to test the application while the source Microsoft SQL Server remains online. After testing, you can schedule application downtime and rerun the PowerShell backup script during your final cutover.

The entire replatforming process can also be automated and run unattended. You can run the Systems Manager SSM document [AWSEC2-SQLServerDBRestore](#) to import your existing database backup files into Microsoft SQL Server on EC2 Linux without using the PowerShell backup script.

# Components

The Windows to Linux replatforming assistant for Microsoft SQL Server Databases script consists of two main components:

1. A [PowerShell backup script](#), which backs up on-premises Microsoft SQL Server databases to an Amazon S3 storage bucket. It then invokes the SSM Automation document [AWSEC2-SQLServerDBRestore](#) to restore the backups to a Microsoft SQL Server on EC2 Linux instance.
2. An SSM Automation document named `AWSEC2-SQLServerDBRestore`, which restores database backups to Microsoft SQL Server on EC2 Linux. This automation restores Microsoft SQL Server database backups stored in Amazon S3 to Microsoft SQL Server 2017 running on an EC2 Linux instance. You can provide your own EC2 instance running Microsoft SQL Server 2017 Linux, or the automation launches and configures a new EC2 instance with Microsoft SQL Server 2017 on Ubuntu 16.04. The automation supports the restoration of full, differential, and transactional log backups, and accepts multiple database backup files. The automation automatically restores the most recent valid backup of each database in the files provided. For more information, see [AWSEC2-SQLServerDBRestore](#).

## Replatforming script prerequisites

This section covers the steps necessary to run the Windows to Linux replatforming script.

### Contents

- [Prerequisites to run the replatforming script](#)
- [Prerequisites for replatforming to an existing EC2 instance](#)

## Prerequisites to run the replatforming script

In order to run the Windows to Linux replatforming assistant for Microsoft SQL Server Databases script, you must do the following:

### 1. Install the AWS PowerShell module

To install the AWS PowerShell module, follow the steps listed in [Installing the AWS Tools for PowerShell on Windows](#). We recommend that you use PowerShell 3.0 or later for the backup script to work properly.

## 2. Install the Windows to Linux replatforming assistant PowerShell backup script

To run the Windows to Linux replatforming assistant, download the PowerShell backup script: [MigrateSQLServerToEC2Linux.ps1](#).

## 3. Add an AWS user profile to the AWS SDK store

To add and configure the AWS user profile, see the steps listed in [Managing Profiles](#) in the *AWS Tools for PowerShell User Guide*. [Set the following IAM policy](#) for your user profile.

JSON

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "VisualEditor0",
      "Effect": "Allow",
      "Action": "iam:PassRole",
      "Resource": "arn:aws:iam::123456789012:role/DevTeam*"
    },
    {
      "Sid": "VisualEditor1",
      "Effect": "Allow",
      "Action": [
        "ec2:RebootInstances",
        "ssm:SendCommand",
        "ssm:GetAutomationExecution",
        "ec2:DescribeInstances",
        "ssm:ListCommands",
        "ec2:CreateTags",
        "s3:CreateBucket",
        "ec2:RunInstances",
        "s3:ListBucket",
        "ssm:GetCommandInvocation",
        "s3:PutEncryptionConfiguration",
        "ec2:DescribeImages",
        "s3:PutObject",
        "s3:GetObject",
        "ssm:StartAutomationExecution",
        "ssm:DescribeInstanceInformation",
        "s3:DeleteObject",
        "ssm:ListCommandInvocations",

```

```
        "s3:DeleteBucket",
        "ec2:DescribeInstanceStatus"
    ],
    "Resource": "*"
  }
]
```

#### 4. Create an IAM instance profile role

To create an IAM instance profile role in order to run Systems Manager on EC2 Linux, see the steps listed under [Create an IAM instance profile for Systems Manager](#) in the *AWS Systems Manager User Guide*.

## Prerequisites for replatforming to an existing EC2 instance

To replatform to an existing instance running Microsoft SQL Server 2017 on Linux, you must:

1. Configure the EC2 instance with an AWS Identity and Access Management (IAM) instance profile and attach the AmazonSSMManagedInstanceCore managed policy.

For information about creating an IAM instance profile for Systems Manager and attaching it to an instance, see the following topics in the *AWS Systems Manager User Guide*:

- [Create an IAM instance profile for Systems Manager](#)
  - [Attach an IAM instance profile to an Amazon EC2 instance](#)
2. Verify that SSM Agent is installed on your EC2 instance. For more information, see [Working with SSM Agent on EC2 instances for Windows Server](#) in the *AWS Systems Manager User Guide*.
  3. Verify that the EC2 instance has enough free disk space to download and restore the Microsoft SQL Server backups.

## Run the Windows to Linux replatforming assistant for SQL Server script

This section contains the PowerShell parameter definitions and scripts for replatforming your databases. For more information about how to use PowerShell scripts, see [PowerShell](#).

### Topics

- [Replatforming script examples](#)
- [Replatforming script parameters](#)

## Replatforming script examples

The following common scenarios and example PowerShell scripts demonstrate how to replatform your Microsoft SQL Server databases using Windows to Linux replatforming assistant for Microsoft SQL Server Databases.

### Important

The Windows to Linux Replatforming Assistant for Microsoft SQL Server Databases resets the SQL Server server administrator (SA) user password on the target instance every time that it is run. After the replatform process is complete, you must set your own SA user password before you can connect to the target SQL Server instance.

### Syntax

The Windows to Linux replatforming assistant for Microsoft SQL Server Databases script adheres to the syntax shown in the following example.

```
PS C:\> C:\MigrateSQLServerToEC2Linux.ps1 [[-SqlServerInstanceName] <String>] [[-DBNames]<Object[]>] [-MigrateAllDBs] [PathForBackup] <String> [-SetSourceDBModeReadOnly] [-IamInstanceProfileName] <String> [-AWSRegion] <String> [[-EC2InstanceId] <String>] [[-EC2InstanceType] <String>] [[-EC2KeyPair] <String>] [[-SubnetId] <String>] [[-AWSProfileName] <String>] [[-AWSProfileLocation] <String>] [-GeneratePresignedUrls] [<CommonParameters>]
```

### Example 1: Move a database to an EC2 instance

The following example shows how to move a database named AdventureDB to an EC2 Microsoft SQL Server on Linux instance, with an instance ID of `i-024689abcdef`, from the Microsoft SQL Server Instance named MSSQLSERVER. The backup directory to be used is `D:\\Backup` and the AWS Region is `us-east-2`.

```
PS C:\> ./MigrateSQLServerToEC2Linux.ps1 - SQLServerInstanceName MSSQLSERVER -  
EC2InstanceId i-  
024689abcdef -DBNames AdventureDB -PathForBackup D:\\Backup -AWSRegion us-east-2 -  
IamInstanceProfileName AmazonSSMManagedInstanceCore
```

### Example 2: Move a database to an EC2 instance using the AWS credentials profile

The following example shows how to move the database in Example 1 using the AWS credentials profile: DBMigration.

```
PS C:\> ./MigrateSQLServerToEC2Linux.ps1 - SQLServerInstanceName MSSQLSERVER -  
EC2InstanceId i-  
024689abcdef -DBNames AdventureDB -PathForBackup D:\\Backup -AWSRegion us-east-2 -  
AWSProfileName  
DBMigration -IamInstanceProfileName AmazonSSMManagedInstanceCore
```

### Example 3: Move a database to a new m5.large type instance

The following example shows how to create an m5.large type EC2 Linux instance in subnet-abc127 using the Key Pair customer-ec2-keypair and then moving AdventureDB and TestDB to the new instance from the database used in Examples 1 and 2.

```
PS C:\> ./MigrateSQLServerToEC2Linux.ps1 -EC2InstanceType m5.large -SubnetId subnet-  
abc127 -EC2KeyPair  
customer-ec2-keypair -DBNames AdventureDB,TestDB -PathForBackup D:\\Backup -  
AWSRegion us-east-2 -  
AWSProfileName DBMigration -IamInstanceProfileName AmazonSSMManagedInstanceCore
```

### Example 4: Move all databases to a new m5.large type instance

The following example shows how to create an m5.large type EC2 Linux instance in subnet-abc127 using the Key Pair customer-ec2-keypair and then migrating all databases to the instance from databases used in Examples 1 and 2.

```
PS C:\> ./MigrateSQLServerToEC2Linux.ps1 -EC2InstanceType m5.large -SubnetId subnet-  
abc127 -EC2KeyPair  
customer-ec2-keypair -MigrateAllDBs -PathForBackup D:\\Backup -AWSRegion us-east-2 -  
AWSProfileName  
DBMigration -IamInstanceProfileName AmazonSSMManagedInstanceCore
```

## Replatforming script parameters

The following parameters are used by the PowerShell script to replatform your Microsoft SQL Server databases.

### **-SqlServerInstanceName**

The name of the Microsoft SQL Server instance to be backed up. If a value for SqlServerInstanceName is not provided, \$env:ComputerName is used by default.

Type: String

Required: No

### **-DBNames**

The names of the databases to be backed up and restored. Specify the names of the databases in a comma-separated list (for example, adventureDB,universityDB). Either the DBNames or MigrateAllDBs parameter is required.

Type: Object

Required: No

### **-MigrateAllDBs**

This switch is disabled by default. If this switch is enabled, the automation migrates all databases except for the system databases (master, msdb, tempdb). Either the DBNames or MigrateAllDBs parameter is required.

Type: SwitchParameter

Required: No

### **-PathForBackup**

The path where the full backup is stored.

Type: String

Required: Yes

### **-SetSourceDBModeReadOnly**

This switch is disabled by default. If this switch is enabled, it makes the database read-only during migration.

Type: SwitchParameter

Required: No

### **-IamInstanceProfileName**

Enter the AWS IAM instance role with permissions to run Systems Manager Automation on your behalf. See [Getting Started with Automation](#) in the *AWS Systems Manager User Guide*.

Type: String

Required: Yes

### **-AWSRegion**

Enter the AWS Region where your Amazon S3 buckets are created to store database backups.

Type: String

Required: Yes

### **-EC2InstanceId**

To restore Microsoft SQL Server databases to an existing EC2 instance running Microsoft SQL Server Linux, enter the instance ID of the instance. Make sure that the EC2 instance already has the AWS Systems Manager SSM Agent installed and running.

Type: String

Required: No

### **-EC2InstanceType**

To restore Microsoft SQL Server databases to a new EC2 Linux instance, enter the instance type of the instance to be launched.

Type: String

Required: No

### **-EC2KeyPair**

To restore Microsoft SQL Server databases to a new EC2 Linux instance, enter the name of the EC2 Key Pair to be used to access the instance. This parameter is recommended if you are creating a new EC2 Linux instance.

Type: String

Required: No

### **-SubnetId**

This parameter is required when creating a new EC2 Linux instance. When creating a new EC2 Linux instance, if SubnetId is not provided, the AWS user default subnet is used to launch the EC2 Linux instance.

Type: String

Required: No

### **-AWSProfileName**

The name of the AWS profile that the automation uses when connecting to AWS services. For more information on the required user permissions, see [Getting Started with Automation](#) in the *AWS Systems Manager User Guide*. If a profile is not entered, the automation uses your default AWS profile.

Type: String

Required: No

### **-AWSProfileLocation**

The location of the AWS Profile if the AWS Profile is not stored in the default location.

Type: String

Required: No

### **-GeneratePresignedUrls**

This parameter is only used when replatforming to non-EC2 instances, such as to VMware Cloud on AWS or on-premises VMs.

Type: SwitchParameter

Required: No

### <CommonParameters>

This cmdlet supports the common parameters: Verbose, Debug, ErrorAction, ErrorVariable, WarningAction, WarningVariable, OutBuffer, PipelineVariable, and OutVariable. For more information, see [About Common Parameters](#) in the Microsoft PowerShell documentation.

Required: No

# Security in Microsoft SQL Server on Amazon EC2

Cloud security at AWS is the highest priority. As an AWS customer, you benefit from a data center and network architecture that is built to meet the requirements of the most security-sensitive organizations.

Security is a shared responsibility between AWS and you. The [shared responsibility model](#) describes this as security of the cloud and security in the cloud:

- **Security of the cloud** – AWS is responsible for protecting the infrastructure that runs AWS services in the AWS Cloud. AWS also provides you with services that you can use securely. Third-party auditors regularly test and verify the effectiveness of our security as part of the [AWS Compliance Programs](#). To learn about the compliance programs that apply to SQL Server on EC2, see [AWS Services in Scope by Compliance Program](#).
- **Security in the cloud** – Your responsibility is determined by the AWS service that you use. You are also responsible for other factors including the sensitivity of your data, your company's requirements, and applicable laws and regulations

For detailed information about how to configure Amazon EC2 to meet your security and compliance objectives, see [Security in Amazon EC2](#) in the *User Guide for Windows Instances*.

## Identity and access management for Microsoft SQL Server on Amazon EC2

AWS Identity and Access Management (IAM) is an AWS service that helps an administrator securely control access to AWS resources. IAM administrators control who can be *authenticated* (signed in) and *authorized* (have permissions) to use SQL Server on EC2 resources. IAM is an AWS service that you can use with no additional charge.

Your security credentials identify you to services in AWS and grant you access to AWS resources, such as your Amazon EC2 resources. You can use features of Amazon EC2 and IAM to allow other users, services, and applications to use your Amazon EC2 resources without sharing your security credentials. You can use IAM to control how other users use resources in your AWS account, and you can use security groups to control access to your Amazon EC2 instances. You can choose to allow full or limited use of your Amazon EC2 resources.

If you are a developer, you can use IAM roles to manage the security credentials needed by the applications that you run on your EC2 instances. After you attach an IAM role to your instance, your applications running on the instance can retrieve the credentials from the Instance Metadata Service (IMDS).

For best practices for securing your AWS resources using IAM, see [Security best practices in IAM](#) in the *IAM User Guide*.

## Contents

- [AWS managed policies for Microsoft SQL Server on Amazon EC2](#)
- [Service-linked role for Microsoft SQL Server on Amazon EC2](#)

## AWS managed policies for Microsoft SQL Server on Amazon EC2

To add permissions to users, groups, and roles, it is easier to use AWS managed policies than to write policies yourself. It takes time and expertise to [create IAM customer managed policies](#) that provide your team with only the permissions they need. To get started quickly, you can use our AWS managed policies. These policies cover common use cases and are available in your AWS account. For more information about AWS managed policies, see [AWS managed policies](#) in the *IAM User Guide*.

AWS services maintain and update AWS managed policies. You can't change the permissions in AWS managed policies. Services occasionally add additional permissions to an AWS managed policy to support new features. This type of update affects all identities (users, groups, and roles) where the policy is attached. Services are most likely to update an AWS managed policy when a new feature is launched or when new operations become available. Services do not remove permissions from an AWS managed policy, so policy updates won't break your existing permissions.

Additionally, AWS supports managed policies for job functions that span multiple services. For example, the **ReadOnlyAccess** AWS managed policy provides read-only access to all AWS services and resources. When a service launches a new feature, AWS adds read-only permissions for new operations and resources. For a list and descriptions of job function policies, see [AWS managed policies for job functions](#) in the *IAM User Guide*.

### AWS managed policy: AWSEC2SqlHaInstancePolicy

You can attach this managed policy to the IAM role that's attached to your Amazon EC2 High Availability for SQL Server instance. The policy grants permissions to execute AWS owned Systems

Manager command document **AWSEC2-DetectSqlHaState** to the instance, to retrieve the EC2 SQL HA instance metadata and decide whether it's in active or standby state.

To view the permissions for this policy, see [AWSEC2SqlHaInstancePolicy](#) in the *AWS Managed Policy Reference*.

## AWS managed policy: AWSEC2SqlHaServiceRolePolicy

This policy is attached to the service-linked role named [AWSServiceRoleForEC2SqlHa](#) to allow Amazon EC2 High Availability for SQL Server on EC2 to detect whether an EC2 instance that's tagged with the EC2 SQL High Availability identifier (SqlHaMonitored set to true) is running in active or standby mode.

To view the permissions for this policy, see [AWSEC2SqlHaServiceRolePolicy](#) in the *AWS Managed Policy Reference*.

## AWS managed policy: AWSEC2VssRestorePolicy

You can attach this managed policy to the IAM role that's used to execute the AWSEC2 - RestoreSqlServerDatabaseWithVss automation runbook. The policy grants permissions to create volumes from VSS snapshots, attach them to instances, and invoke AWS Systems Manager Run Command documents required for database restoration.

### Permissions details

This policy includes the following permissions:

- **ec2** – Allows principals to create volumes from snapshots tagged with AwsVssConfig, attach volumes to instances, tag volumes during creation, and describe volumes, snapshots, and instance attributes.
- **ssm** – Allows principals to describe SSM managed instances, retrieve SSM Run Command documents required for VSS restore operations, send commands to instances, and list command invocations and executions.

To view the permissions for this policy, see [AWSEC2VssRestorePolicy](#) in the *AWS Managed Policy Reference*.

## SQL Server on EC2 updates to AWS managed policies

View details about updates to AWS managed policies for SQL Server on EC2 since this service began tracking these changes.

| Change                                                    | Description                                                                                                                                                                                                                                                                            | Date              |
|-----------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|
| <a href="#">AWSEC2VssRestorePolicy</a> – New policy       | Added the <b>AWSEC2VssRestorePolicy</b> policy that can be attached to the IAM role assumed by the <code>AWSEC2-RestoreSqlS</code> <code>erverDatabaseWithVss</code> automation runbook for restoring Microsoft SQL Server databases from VSS snapshots.                               | March 25, 2026    |
| <a href="#">AWSEC2SqlHaInstancePolicy</a> – New policy    | Added the <b>AWSEC2SqlHaInstancePolicy</b> policy that can be attached to IAM role that's attached to the Windows and SQL HA instance to facilitate metadata collection for the purpose of keeping track of the current state of the database as it applies to active or passive mode. | November 17, 2025 |
| <a href="#">AWSEC2SqlHaServiceRolePolicy</a> – New policy | Added the policy that's attached to the <b>AWSServiceRoleForEC2SqlHa</b> service-linked role to detect whether an EC2 instance that's tagged with the EC2 SQL High Availability identifier is                                                                                          | November 17, 2025 |

| Change                                     | Description                                                            | Date              |
|--------------------------------------------|------------------------------------------------------------------------|-------------------|
|                                            | running in standby or passive mode.                                    |                   |
| SQL Server on EC2 started tracking changes | SQL Server on EC2 started tracking changes to its AWS managed policies | November 17, 2025 |

## Service-linked role for Microsoft SQL Server on Amazon EC2

Amazon EC2 uses service-linked roles for the permissions that it requires to call other AWS services on your behalf. A service-linked role is a unique type of IAM role that is linked directly to an AWS service. Service-linked roles provide a secure way to delegate permissions to AWS services because only the linked service can assume a service-linked role. For more information about how Amazon EC2 uses IAM roles, including service-linked roles, see [IAM roles for Amazon EC2](#) in the *Amazon EC2 User Guide*.

### Service-linked role permissions for Amazon EC2 High Availability for SQL Server

Amazon EC2 High Availability for SQL Server uses the service-linked role named **AWSServiceRoleForEC2SqlHa** to allow the service to detect whether an EC2 instance that's tagged with the EC2 SQL High Availability identifier (SqlHaMonitored set to true) is running in active or passive mode.

The AWSServiceRoleForEC2SqlHa service-linked role trusts the following service to assume the role: `ec2sqlha.amazonaws.com`

Amazon EC2 uses the [AWSEC2SqlHaServiceRolePolicy](#) managed policy to complete the following actions:

- **Amazon EC2** – Access is granted for the EC2 SQL High Availability service to describe EC2 instances, instance attributes, instance status which are tagged with the service identifier (SqlHaMonitored set to true).
- **Amazon EventBridge** – Includes access to create Amazon EventBridge event rules and retrieve details about or delete rules that it created. This is to allow the System Manager document **AWSEC2-DetectSqlHaState** execution output being forwarded to the service. A managed Amazon EventBridge rule will be created to forward System Manager run command events.

Managed rules are predefined by User Notifications and include event patterns that are required by the service to manage customer notifications, and unless defined otherwise, only the owning service can utilize these managed rules.

- **AWS Systems Manager** – Includes access to describe instance information and list commands and command invocations. To run the command document that begins with **AWSEC2-DetectSqlHaState**, on a monitored instance, access is granted for the `SendCommand` and `GetCommandInvocation` operations to EC2 SQL Server instances tagged with the service identifier(`SqlHaMonitored` set to `true`).

To view the permissions for this policy, see [AWSEC2SqlHaServiceRolePolicy](#) in the *AWS Managed Policy Reference*.

For more information about using managed policies for EC2 instances, see [AWS managed policies for Amazon EC2](#) in the *Amazon EC2 User Guide*.

# Document history for the Microsoft SQL Server on Amazon EC2 User Guide

The following table describes the documentation releases for Microsoft SQL Server on Amazon EC2.

| Change                                                                            | Description                                                                                                                                                          | Date              |
|-----------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|
| <a href="#">EC2 High Availability for SQL Server</a>                              | Amazon EC2 High Availability for SQL Server allows you to configure instances running license-included SQL Server as part of a HA cluster to reduce licensing costs. | November 17, 2025 |
| <a href="#">Database backup and restore with AWS VSS solution based snapshots</a> | Added support for AWS VSS solution based backup and restore feature.                                                                                                 | January 17, 2025  |
| <a href="#">Downgrade your SQL Server edition</a>                                 | Added a new section about downgrading your Microsoft SQL Server edition.                                                                                             | August 28, 2023   |
| <a href="#">Migration</a>                                                         | Added a new section about migrating to Microsoft SQL Server on Amazon EC2.                                                                                           | August 1, 2023    |
| <a href="#">Initial release</a>                                                   | Initial release of the Microsoft SQL Server on Amazon EC2 User Guide                                                                                                 | August 18, 2022   |